

Adversarial Defense

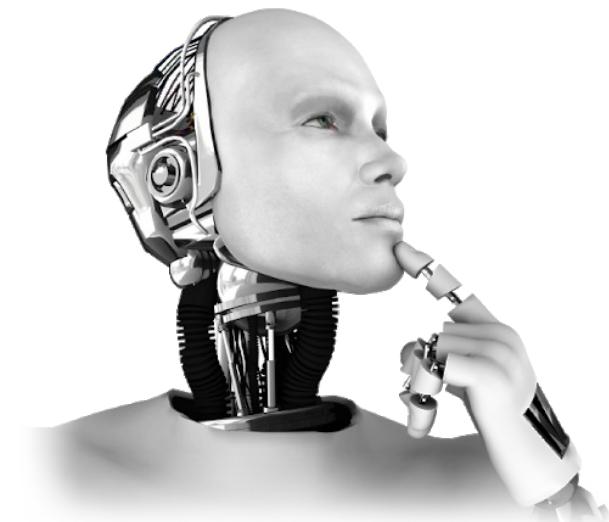
马兴军，复旦大学 计算机学院



Recap: week 4

1. Adversarial Example Detection

- Secondary Classification Methods (二级分类法)
- Principle Component Analysis (主成分分析法, PCA)
- Distribution Detection Methods (分布检测法)
- Prediction Inconsistency (预测不一致性)
- Reconstruction Inconsistency (重建不一致性)
- Trapping Based Detection (诱捕检测法)



Adversarial Attack Competition

RESULTS							
#	User	Entries	Date of Last Entry	Score ▲	Efficiency Score ▲	Error Rate ▲	Detailed Results
1	abcdhhhh	8	10/10/23	0.7042 (3)	0.9901 (1)	0.4183 (5)	View
1	alex_z	2	10/10/23	0.7004 (4)	0.9901 (1)	0.4108 (6)	View
1	siyuandu	2	10/09/23	0.7004 (4)	0.9901 (1)	0.4108 (6)	View
1	xieyong	3	10/06/23	0.7004 (4)	0.9901 (1)	0.4108 (6)	View
1	YiY	1	10/02/23	0.7004 (4)	0.9901 (1)	0.4108 (6)	View
1	jxzhou	7	10/09/23	0.5607 (8)	0.9901 (1)	0.1314 (8)	View
2	wnllixiao	4	10/11/23	0.7077 (1)	0.9802 (2)	0.4353 (4)	View
2	tdlh1	10	10/09/23	0.7077 (1)	0.9802 (2)	0.4353 (4)	View
2	starch	4	10/08/23	0.7077 (1)	0.9802 (2)	0.4353 (4)	View
2	archen	7	10/09/23	0.6516 (6)	0.9802 (2)	0.3231 (7)	View
3	shuyang_jiang	5	10/09/23	0.7069 (2)	0.9703 (3)	0.4436 (3)	View
4	LiGuanyu	2	10/11/23	0.6779 (5)	0.9010 (4)	0.4548 (2)	View
4	hanxunh	1	10/01/23	0.6779 (5)	0.9010 (4)	0.4548 (2)	View
5	X.RW	11	10/08/23	0.6252 (7)	0.7540 (5)	0.4964 (1)	View

Link: https://codalab.lisn.upsaclay.fr/competitions/15669?secret_key=77cb8986-d5bd-4009-82f0-7dde2e819ff8



Adversarial Defense vs Detection

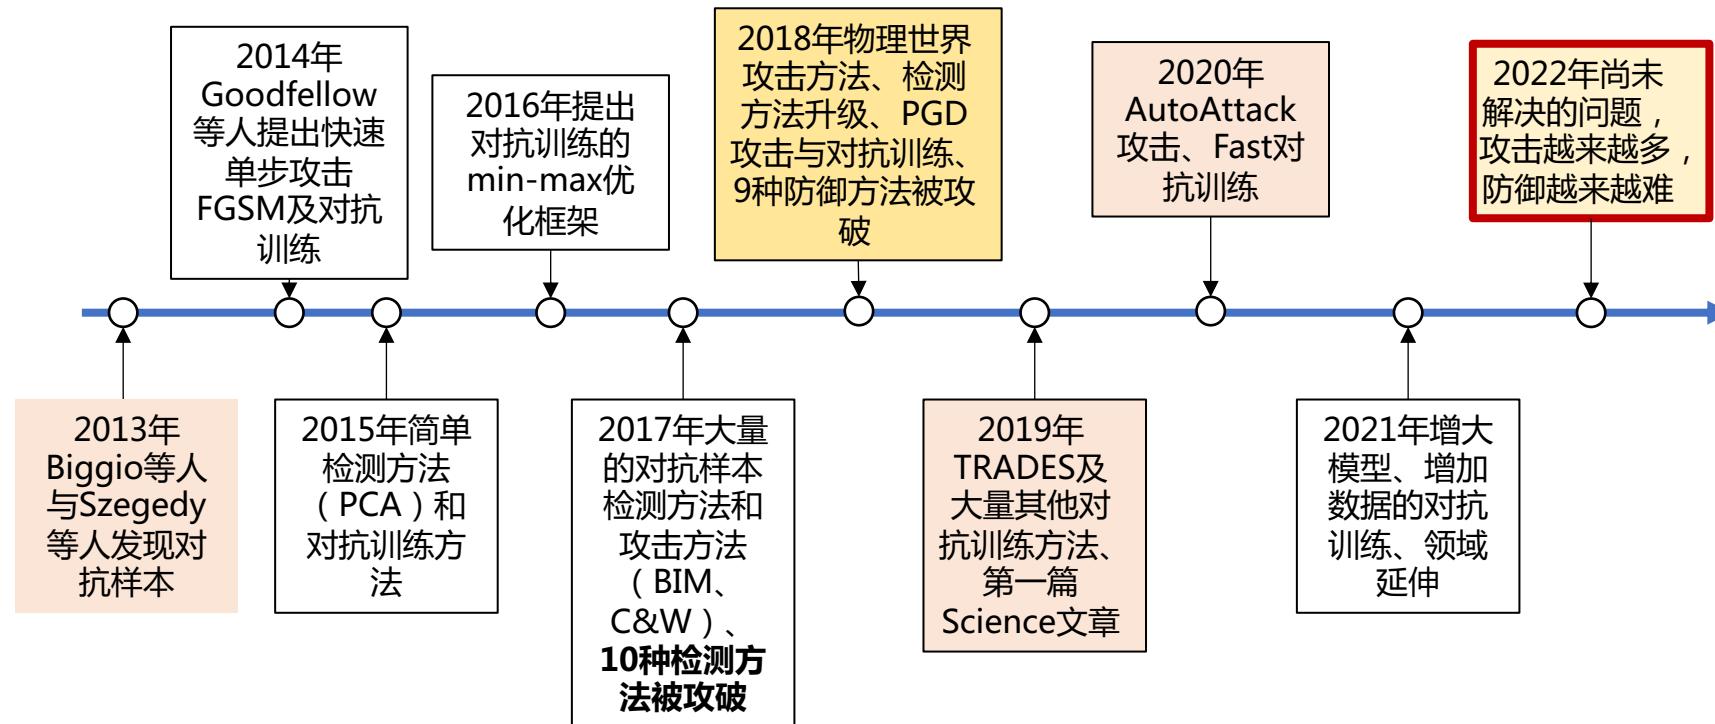
- The weird relationship between defense and detection
 - ✓ Detection **IS** defense
 - ✓ But... when we say **defense**, we (most of the time) mean **the model is secured**, yet detection cannot do that...
 - ✓ In survey papers: detection is defense
 - ✓ In technical papers: defense is defense not detection

- Differences
 - ✓ Defense is to secure the model or the system
 - ✓ Detection is to identify potential threats, which should be followed by a defense strategy, e.g., query rejection (but mostly ignored)
 - ✓ By defense, it mostly means robust training methods

Defense Methods

- Early Defense Methods
- Early Adversarial Training Methods
- Later Adversarial Training Methods
- Remaining Challenges and Recent Progress

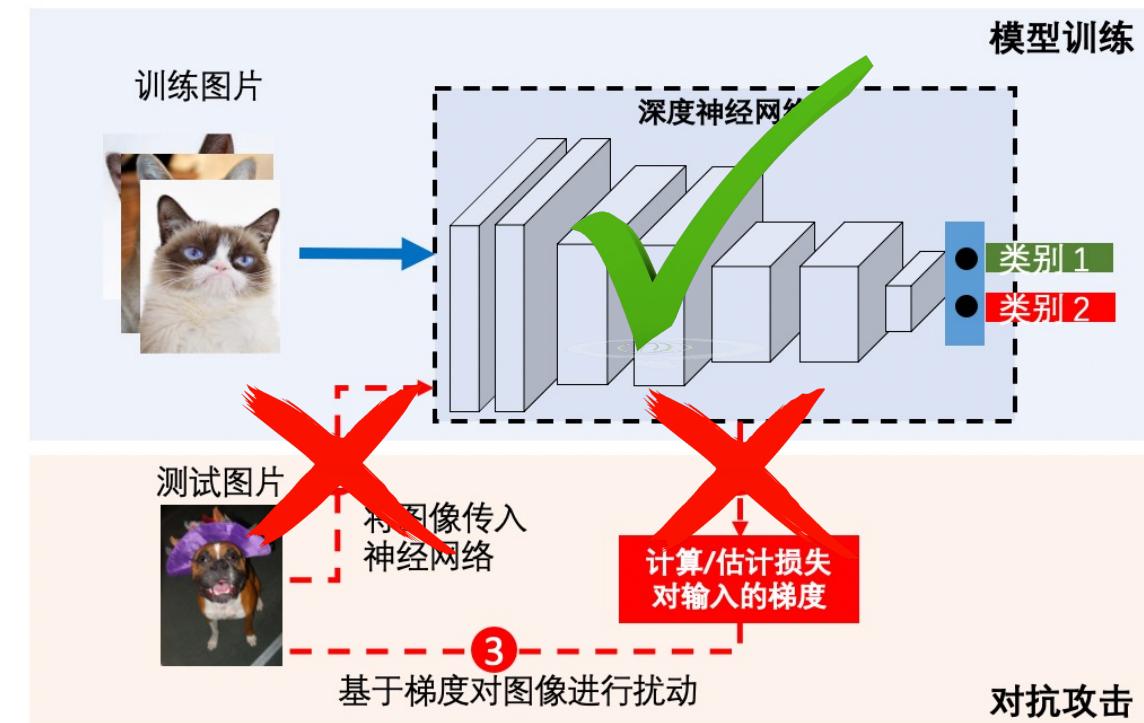
A Recap of the Timeline



Principles of Defense

- Block the attack (**掐头去尾**)
 - Mask the input gradients
 - Regularize the input gradients
 - Distill the logits
 - Denoise the input

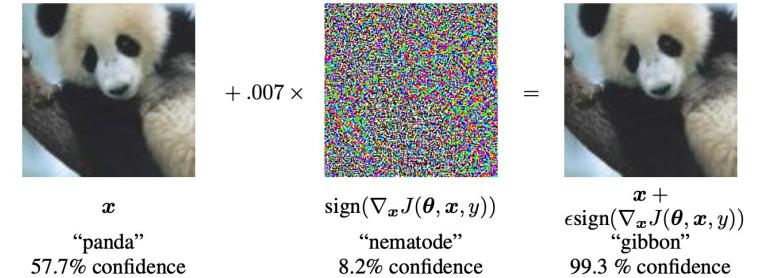
- Robustify the model (**增强中间**)
 - Smooth the decision boundary
 - Reduce the Lipschitzness of the model
 - Smooth the loss landscape



Adversarial Attack

模型训练:

$$\min_{\theta} \sum_{(x_i, y_i) \in D_{train}} L(f_{\theta}(x_i), y_i)$$



对抗攻击:

$$\max_{x'} L(f_{\theta}(x'), y) \text{ subject to } \|x' - x\|_p \leq \epsilon \text{ for } x \in D_{test}$$

分类错误

扰动很小

测试阶段攻击

扰动上限: $\|x' - x\|_{p=1, 2 \text{ or } \infty}$, e.g., $\|\cdot\|_{\infty} \leq \frac{8}{255}$

Performance Metrics

- Measurement of clean performance:

$$\text{accuracy (clean accuracy)} = \frac{\text{correctly classified \textbf{clean} samples}}{\text{total \textbf{clean} samples}}$$

- Measurement of adversarial robustness:

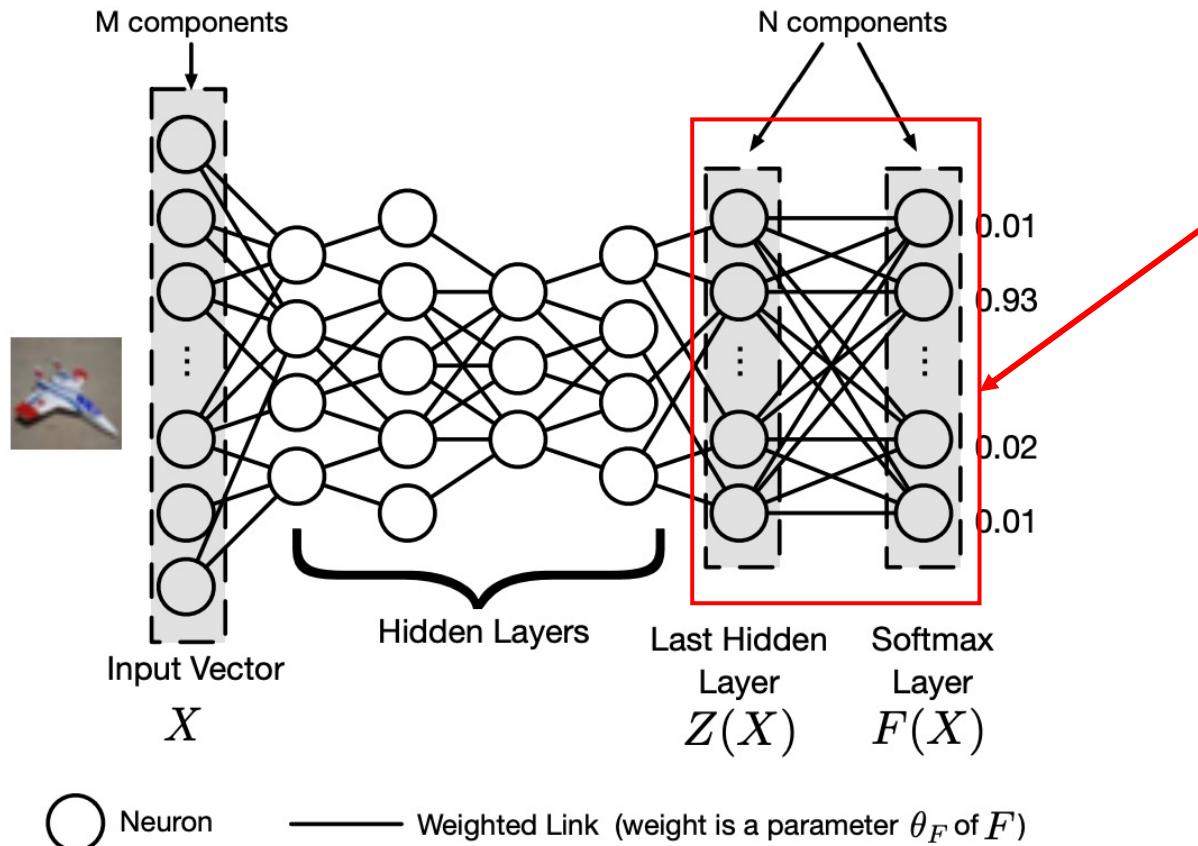
$$\text{robustness (robust accuracy)} = \frac{\text{correctly classified \textbf{advs} samples}}{\text{total \textbf{advs} samples}}$$

- Other metrics: maximum perturbation for 100% attack success rate

Defense Methods

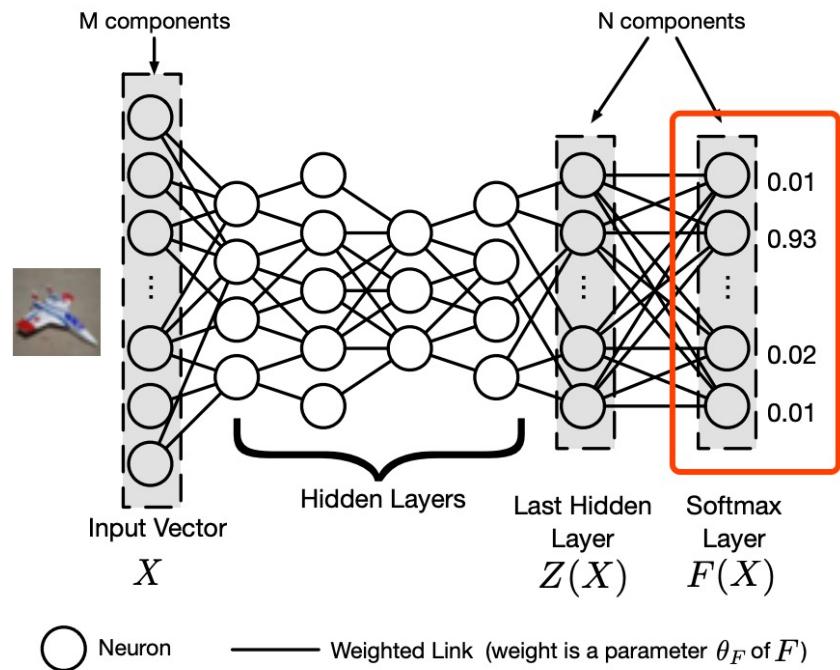
- Early Defense Methods
- Early Adversarial Training Methods
- Advanced Adversarial Training Methods
- Remaining Challenges and Recent Progress

Defensive Distillation



- **Making large logits change to be "small"**
 - Scaling up logits by a few magnitudes;
 - Retrain the last layer with scaled logits;

Defensive Distillation



$$f_{\theta}^i(x) = \frac{e^{z_i(x)}}{\sum_{j=1}^K e^{z_j(x)}}$$

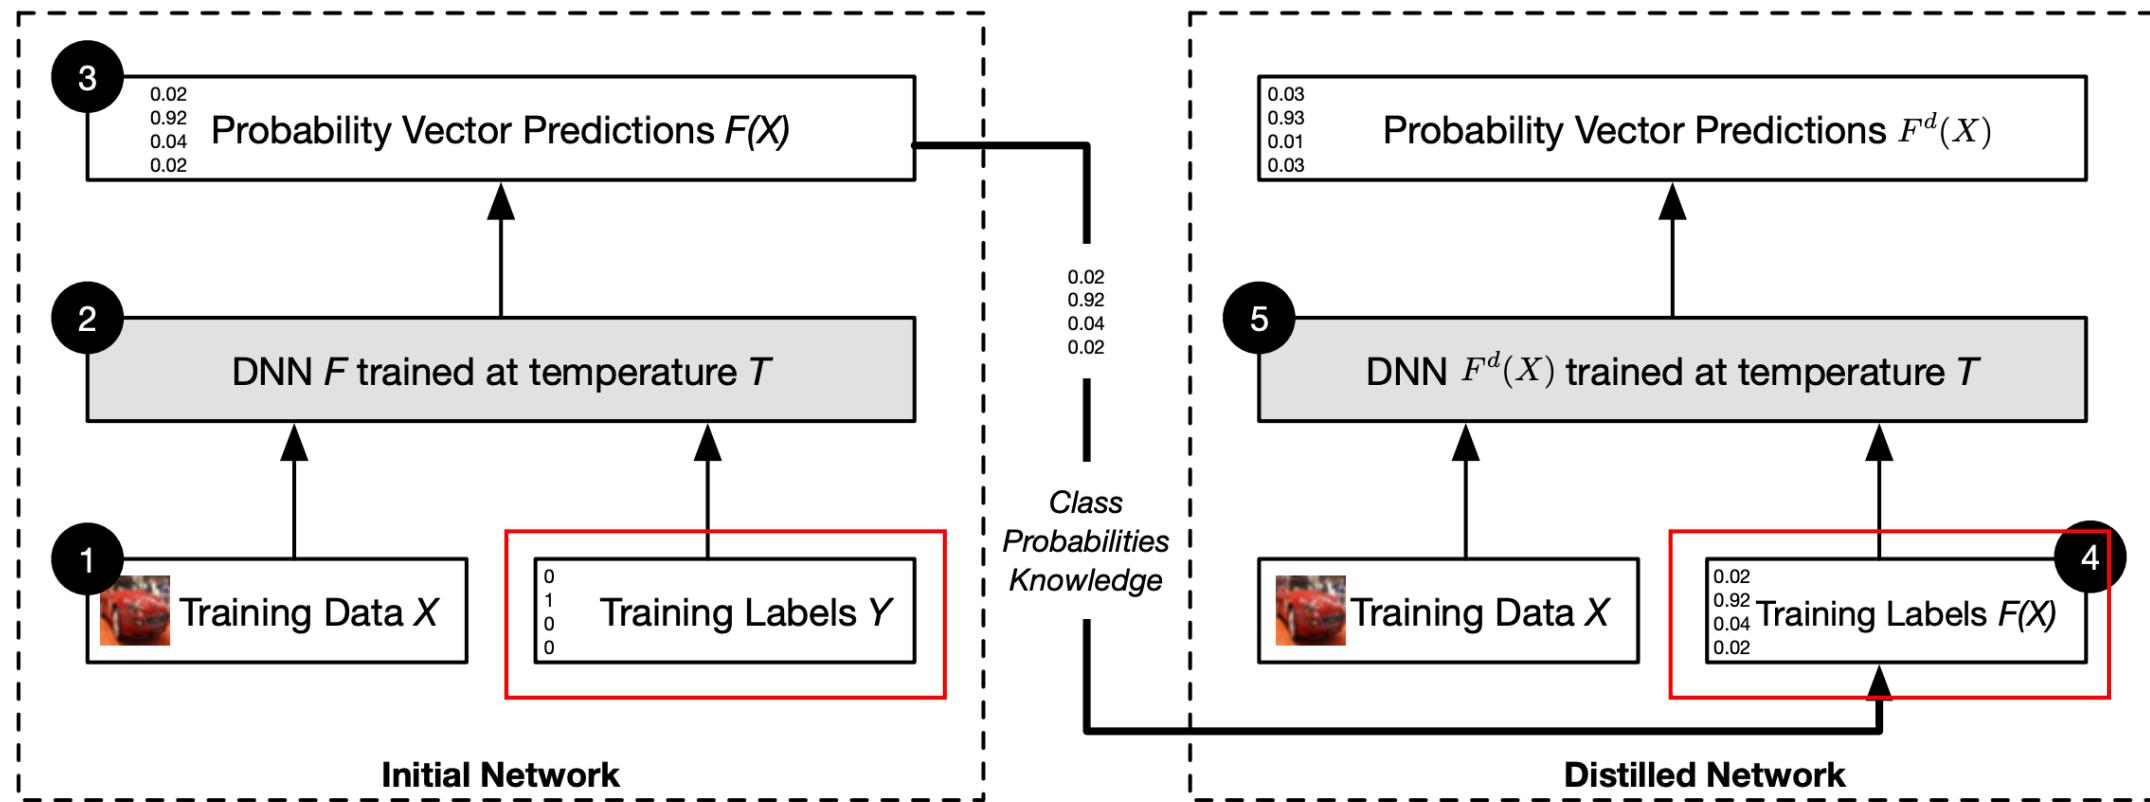
1

Distillation with
temperature T

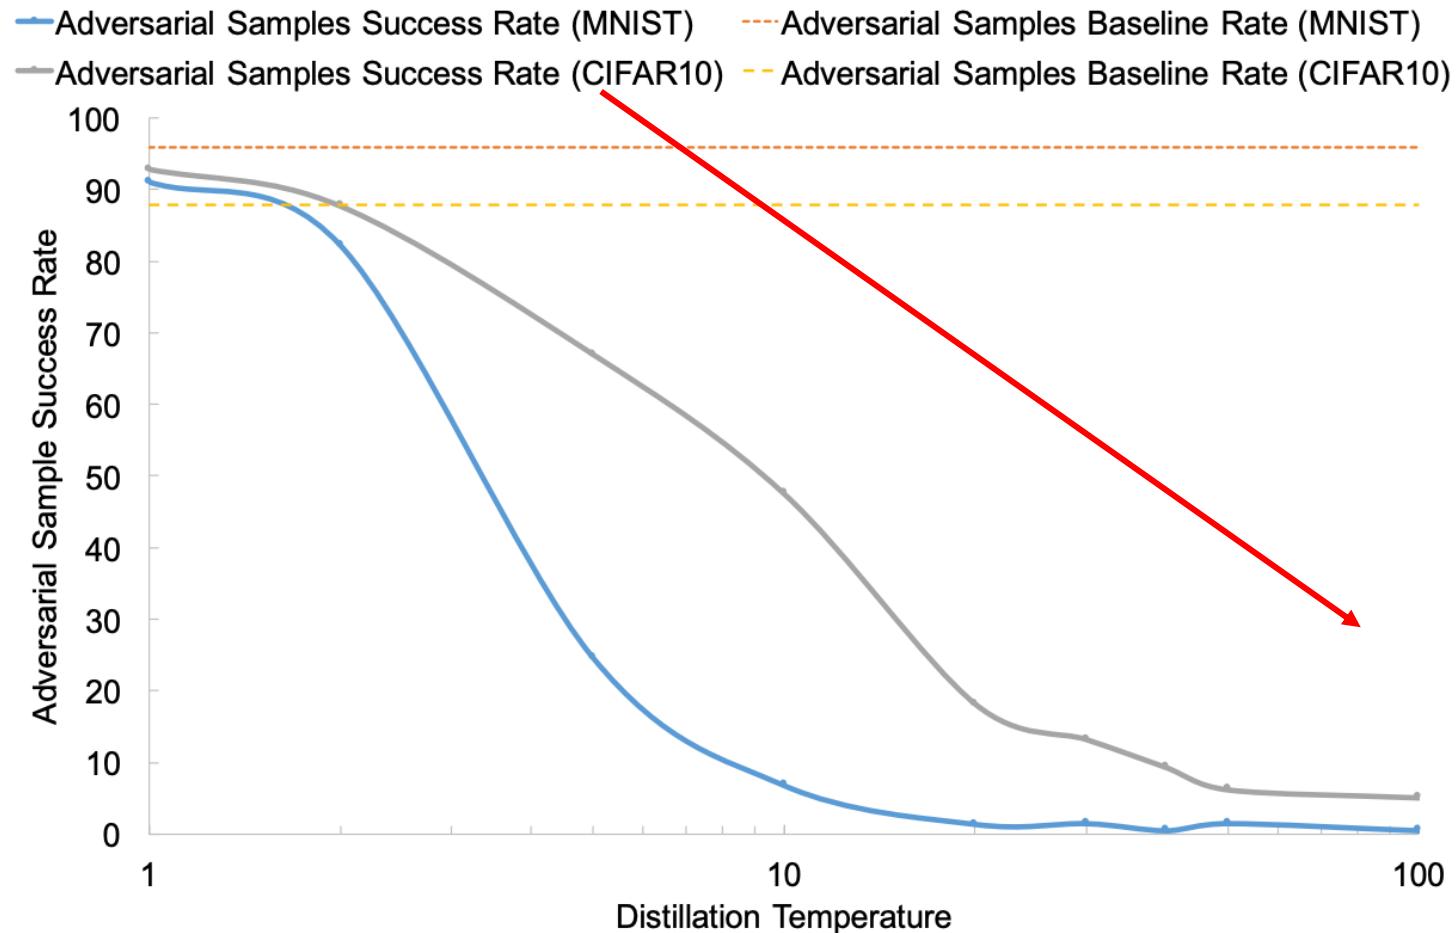
$$f_{\theta}^i(x) = \frac{e^{z_i(x)/T}}{\sum_{j=1}^K e^{z_j(x)/T}}$$

2

Defensive Distillation

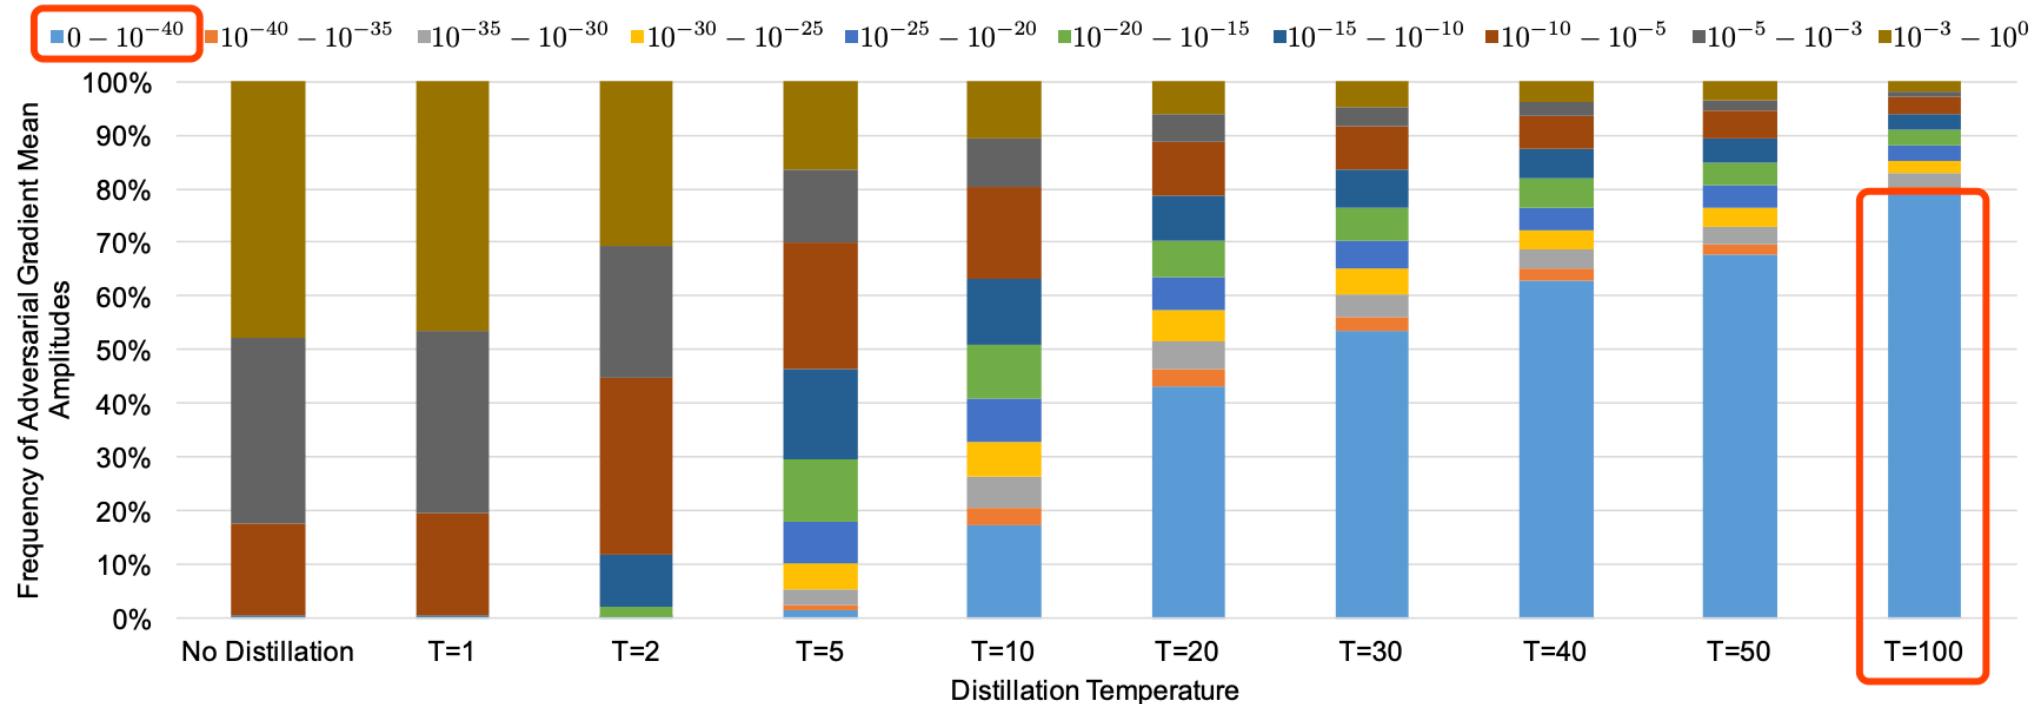


Defensive Distillation



Defensive Distillation

Distillation makes input gradients $\nabla_x L(f_\theta(x), y)$ to be small!

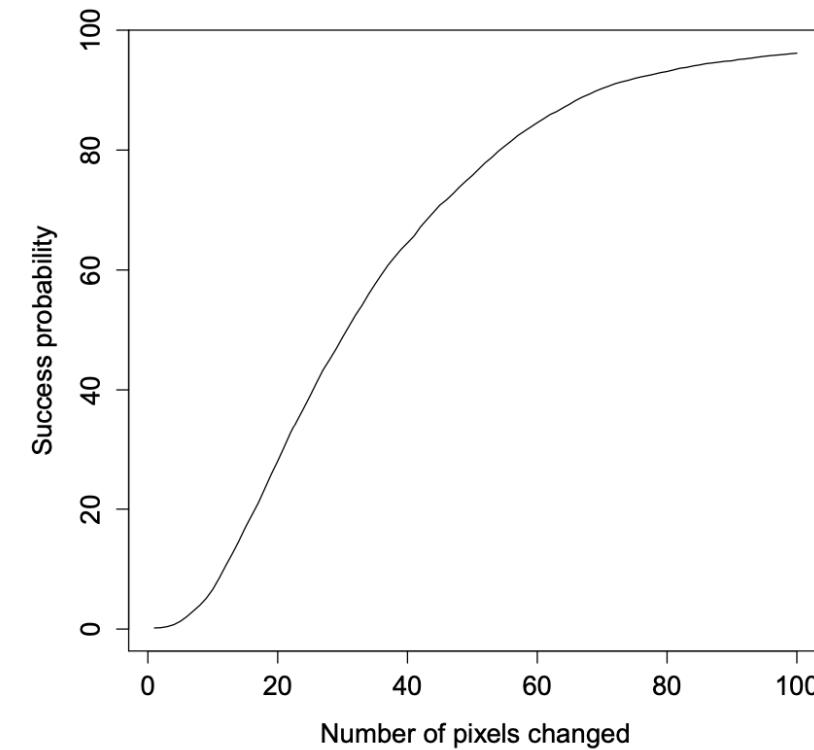


Defensive Distillation Is Not Robust

It can be evaded by attacking the distilled network with the temperature T.

$$x' = x + \varepsilon \cdot \text{sign} \nabla_x L(\hat{f}_\theta(x), y)$$

$$\hat{f}_\theta(x) = \text{softmax}(z(x)/T)$$



Lessons Learned

- Distillation is not a good solution for adversarial robustness
- Vanishing input gradients can still be recovered by a reverse operation
- A defense should be evaluated against the adaptive attack to prove real robustness



Input Gradients Regularization

Directly regularize the input gradients $\nabla_x L(f_\theta(x), y)$ to be small

$$L_{reg} = \underbrace{L(f_\theta(x), y)}_{\text{Classification loss}} + \lambda \underbrace{\|\nabla_x L(f_\theta(x), y)\|_2^2}_{\text{Input gradients regularization}}$$

Related to the **double backpropagation** proposed by Drucker and Le Cun (1992):

$$\arg \min_{\theta} H(y, \hat{y}) + \lambda \|\nabla_x H(y, \hat{y})\|_2^2$$

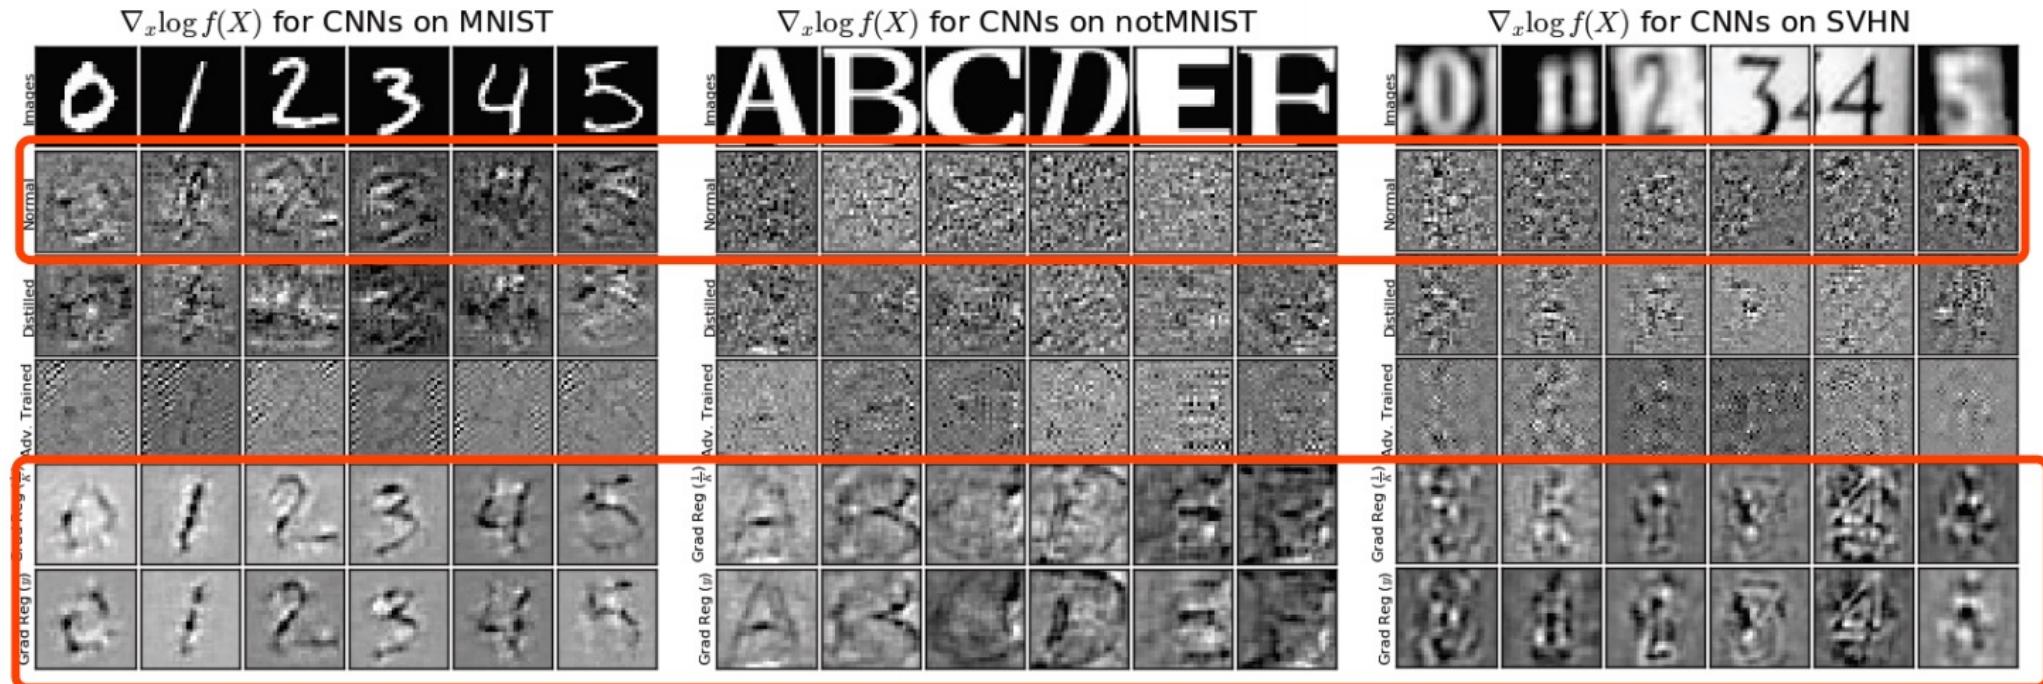


Input Gradients Regularization

蒸馏的

对抗训练的

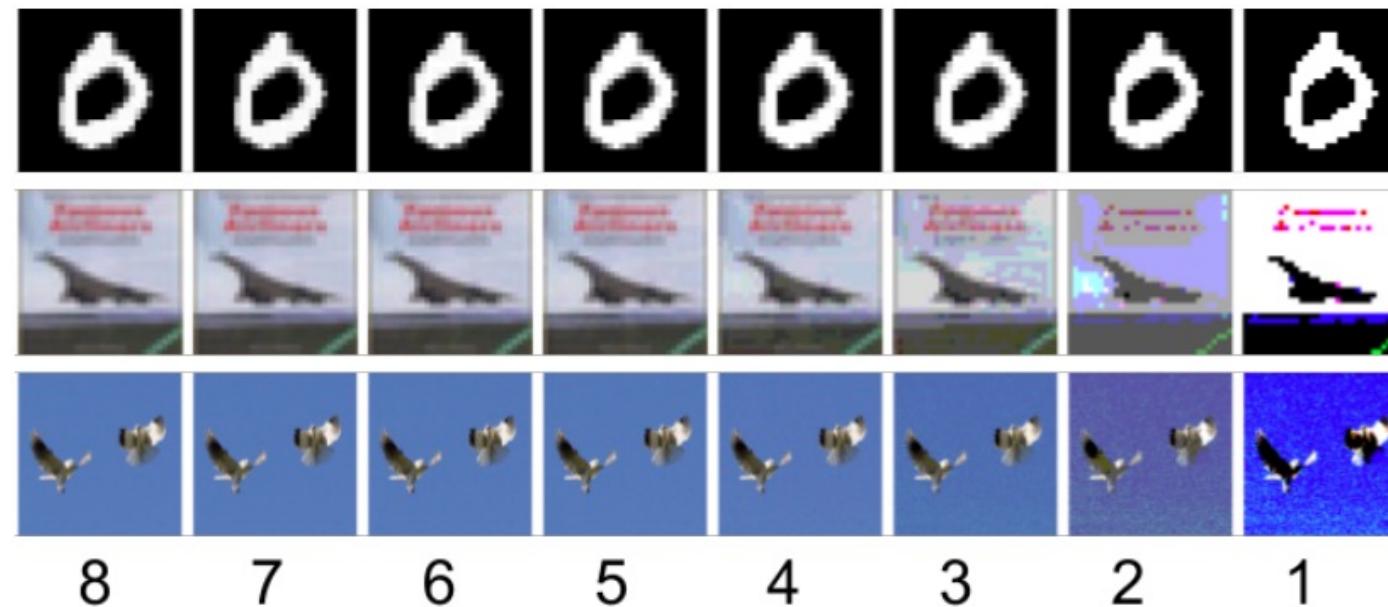
正则化的



Issues: 1) limited adversarial robustness, 2) hurts learning

Feature Squeezing

Compress the input space

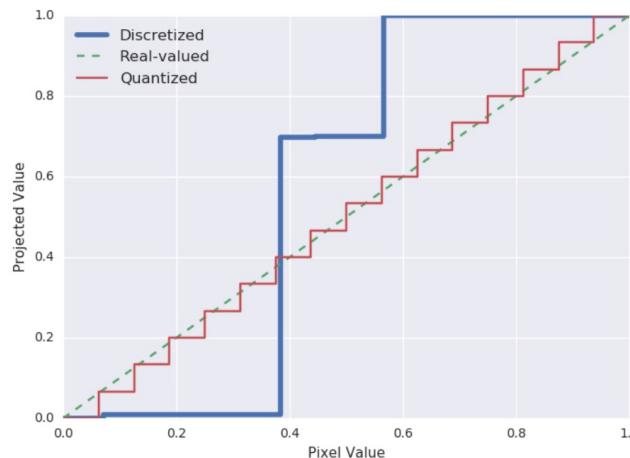


It also hurts performance on large-scale image datasets.

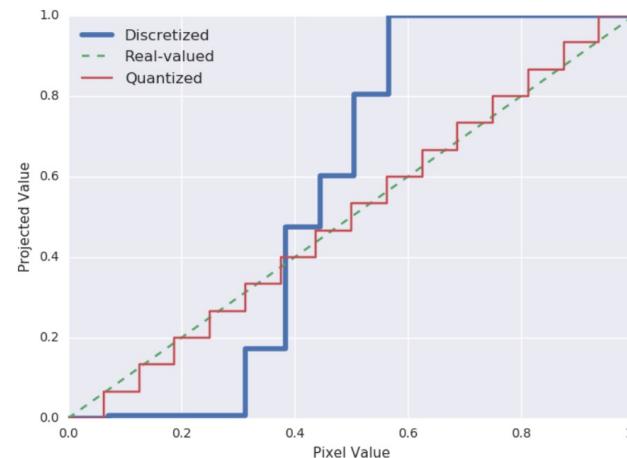
Thermometer Encoding

Discretize the input to break small noise

Real-valued	Quantized	Discretized (one-hot)	Discretized (thermometer)
0.13	0.15	[0100000000]	[0111111111]
0.66	0.65	[0000001000]	[0000001111]
0.92	0.95	[0000000001]	[0000000001]



(a) One hot encoding



(b) Thermometer encoding

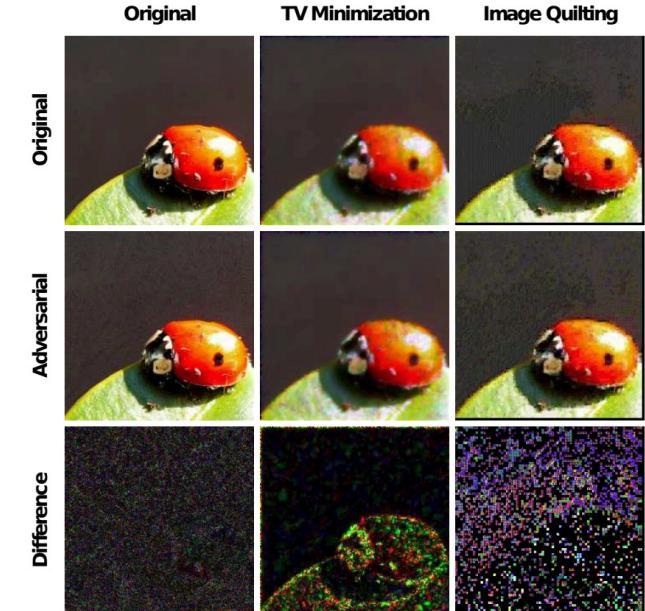
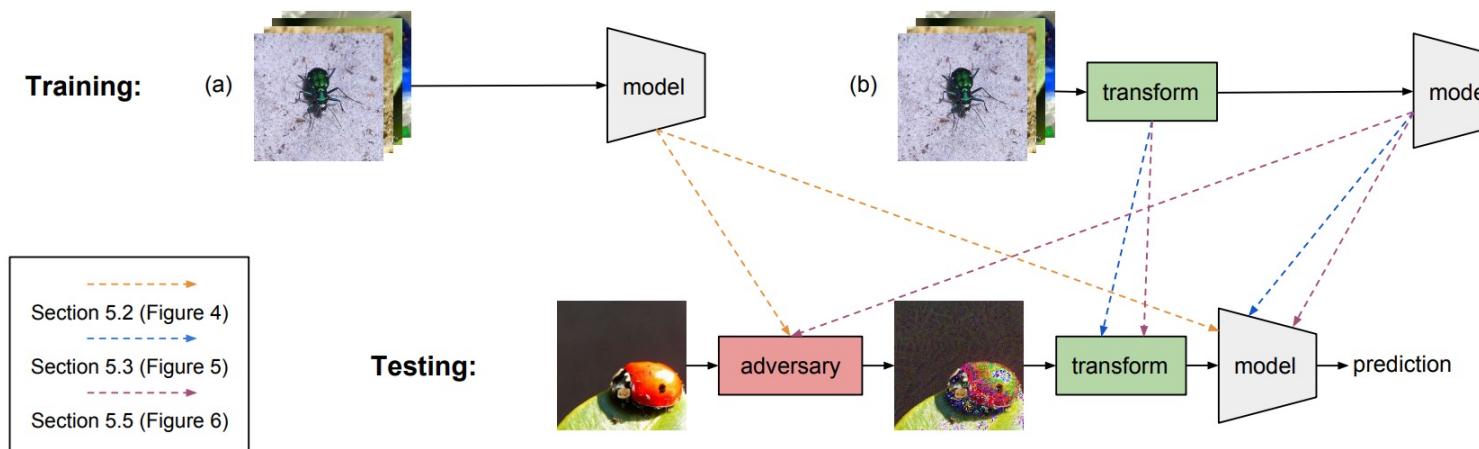
$$\tau(x_{i,j,c})_k = 1 \text{ if } x_{i,j,c} > k/l,$$

$$\tau(0.66) = 1111110000$$

Proposed Thermometer Encoding

Input Transformations

- Image cropping and rescaling
- Bit-depth reduction
- JPEG compression
- Total variance minimization
- Image quilting



Obfuscated Gradients = Fake Robustness

Backward Pass Differentiable Approximation (BPDA): can break **non-differentiable operation** based defenses

$$g(x) \approx f^i(x)$$

find a linear approximation of the non-differentiable operations, e.g., discretization, compression etc.

Expectation Over Transformation (EOT) can break **randomization** based defenses

$$\arg \max_{x'} \mathbb{E}_{t \sim T} [\log P(y_t | t(x'))]$$

T: a set of randomized transformations

$$\text{subject to } \mathbb{E}_{t \sim T} [d(t(x'), t(x))] < \epsilon$$

$$x \in [0, 1]^d$$



BPDA+EOT breaks 7 defenses published at ICLR 2018

Defense	Dataset	Distance	Accuracy
Buckman et al. (2018)	CIFAR	0.031 (ℓ_∞)	0%*
Ma et al. (2018)	CIFAR	0.031 (ℓ_∞)	5%
Guo et al. (2018)	ImageNet	0.005 (ℓ_2)	0%*
Dhillon et al. (2018)	CIFAR	0.031 (ℓ_∞)	0%
Xie et al. (2018)	ImageNet	0.031 (ℓ_∞)	0%*
Song et al. (2018)	CIFAR	0.031 (ℓ_∞)	9%*
Samangouei et al. (2018)	MNIST	0.005 (ℓ_2)	55%**
Madry et al. (2018)	CIFAR	0.031 (ℓ_∞)	47%
Na et al. (2018)	CIFAR	0.015 (ℓ_∞)	15%

We got a survivor!



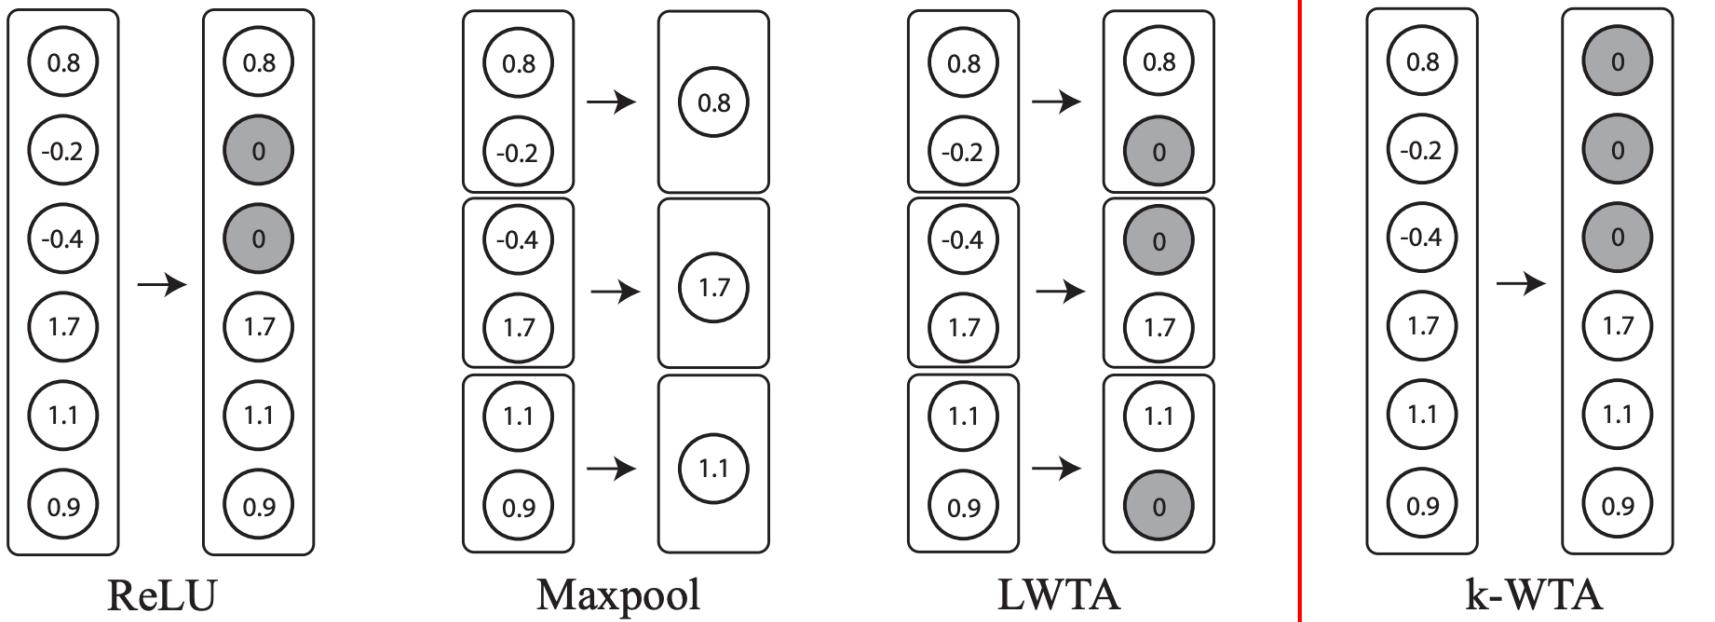
How to Properly Evaluate a Defense?

- ✓ Do not blindly apply multiple (similar) attacks
- ✓ Try at least one gradient-free attack and one hard-label attack
- ✓ Perform a transferability attack using a similar substitute model.
- ✓ For randomized defenses, properly ensemble over randomness
- ✓ For non-differentiable components, apply differentiable techniques (BPDA)
- ✓ Verify that the attacks have converged under the selected hyperparameters
- ✓ Carefully investigate attack hyperparameters and report those selected
- ✓ Compare against prior work and explain important differences
- ✓ Test broader threat models when proposing general defenses



Robust Activation Functions

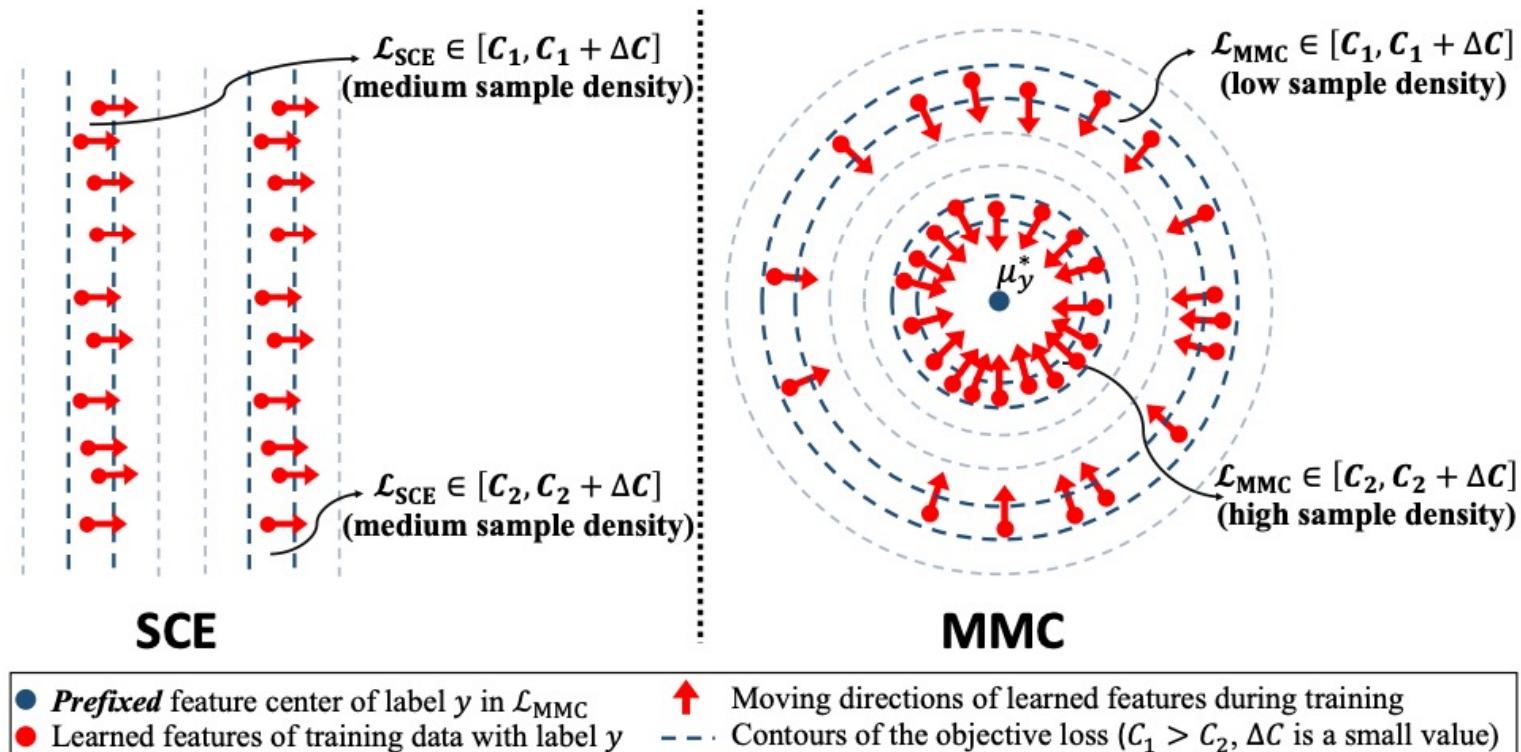
Block the internal activation: break the continuity



k-Winners-Take-All (k-WTA) activation

Robust Loss Function

Max-Mahalanobis center (MMC) loss



Max-Mahalanobis center (MMC); SCE: softmax cross entropy

Robust Inference

Mixup Inference (MI)

$$\hat{x}_k = \lambda x + (1 - \lambda)s_k$$

Algorithm 1 Mixup Inference (MI)

Input: The mixup-trained classifier F ; the input x .

Hyperparameters: The sample distribution p_s ; the mixup ratio λ ; the number of execution N .

Initialize $F_{\text{MI}}(x) = \mathbf{0}$;

for $k = 1$ **to** N **do**

 Sample $y_{s,k} \sim p_s(y_s)$, $x_{s,k} \sim p_s(x_s | y_{s,k})$;

 Mixup x with $x_{s,k}$ as $\tilde{x}_k = \lambda x + (1 - \lambda)x_{s,k}$;

 Update $F_{\text{MI}}(x) = F_{\text{MI}}(x) + \frac{1}{N}F(\tilde{x}_k)$;

end for

Return: The prediction $F_{\text{MI}}(x)$ of input x .

New Adaptive Attacks Break These Defenses

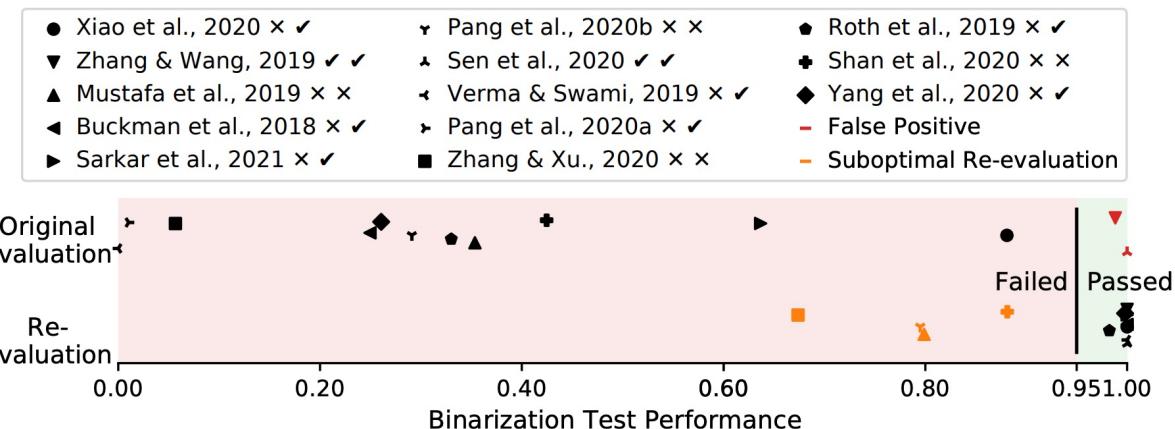
Defense		Attack Themes					
		T1	T2	T3	T4	T5	T6
Appendix B	<i>k</i> -Winners Take All [XZZ20]	●				●	●
Appendix C	<i>The Odds are Odd</i> [RKH19]			●	●		
Appendix D	<i>Generative Classifiers</i> [LBS19]		●	●			
Appendix E	<i>Sparse Fourier Transform</i> [BMV18]	●	●				
Appendix F	<i>Rethinking Cross Entropy</i> [PXD ⁺ 20]			●		●	
Appendix G	<i>Error Correcting Codes</i> [VS19]	●	●				
Appendix H	<i>Ensemble Diversity</i> [PXD ⁺ 19]				●		
Appendix I	<i>EMPIR</i> [SRR20]		●			●	
Appendix J	<i>Temporal Dependency</i> [YLCS19]	●		●	●	●	
Appendix K	<i>Mixup Inference</i> [PXZ20]	●					
Appendix L	<i>ME-Net</i> [YZKX19]	●	●	●			
Appendix M	<i>Asymmetrical Adv. Training</i> [YKR20]			●	●		●
Appendix N	<i>Weakness into a Strength</i> [HYG ⁺ 19]	●	●	●	●		

- T1:** Attack the full defense
- T2:** Target important defense parts
- T3:** Simplify the attack
- T4:** Ensure consistent loss function
- T5:** Optimize with different methods
- T6:** Use strong adaptive attacks

How to Evaluate a Defense?

- Strong attacks:
 - ✓ **AutoAttack** (one must-to-test attack)
 - ✓ **Margin Decomposition (MD) attack** (better than AutoAttack on ViT)
 - ✓ **Minimum-Margin (MM) attack** (new SOTA attack to test?)

- Extra robustness tests
 - ✓ Attack unit tests (Zimmermann et al, 2022)



Croce and Hein. "Reliable evaluation of adversarial robustness with an ensemble of diverse parameter-free attacks." *ICML*, 2020.

Gao et al. Fast and Reliable Evaluation of Adversarial Robustness with Minimum-Margin Attack, ICML 2022.

Zimmermann et al. "Increasing Confidence in Adversarial Robustness Evaluations." *arXiv preprint arXiv:2206.13991* (2022).

Adversarial Training

The idea is simple: just train on adversarial examples!

□ 对抗训练是一种数据增广方法

- 原始数据-> 对抗攻击-> 对抗样本->模型训练

□ 对抗训练是一个min-max鲁棒优化框架:

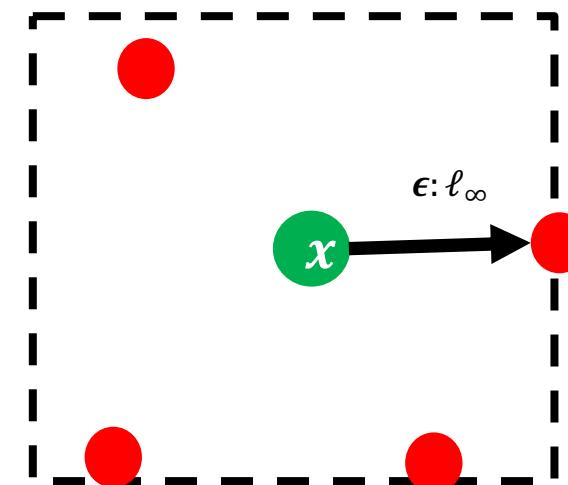
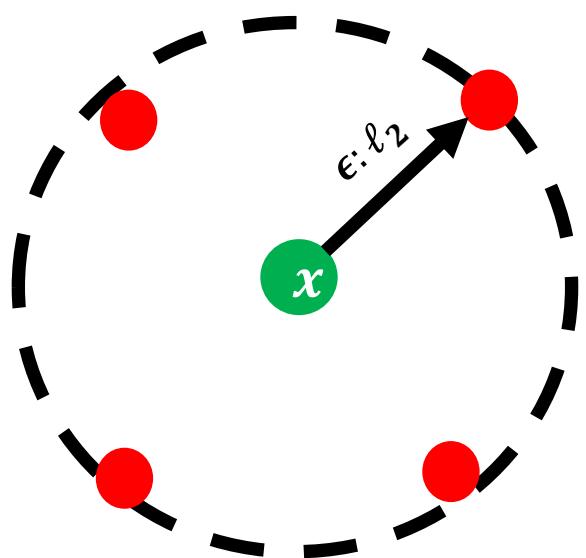
$$\min_{\theta} \frac{1}{n} \sum_{i=1}^n \max_{\|x_i - x_i^0\| \leq \epsilon} L(f_{\theta}(x_i), y_i)$$

$L(f_{\theta}(x_i), y_i) = -y_i \log f_{\theta}(x_i)$ (交叉熵损失函数)

x_i^0 : 原始训练样本, y_i : x_i^0 的正确类别.

Adversarial Training

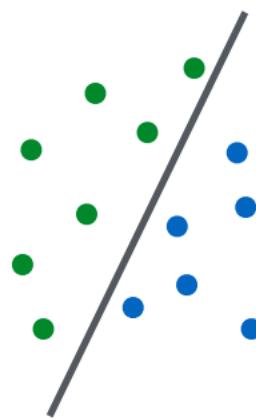
在以 x 为中心的 ϵ -球范围内寻找对抗性最强的样本



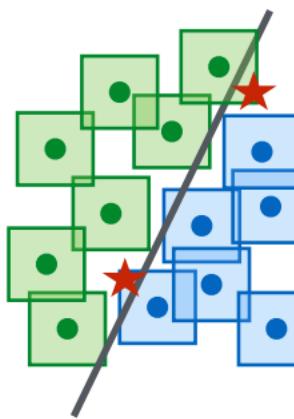
ℓ_∞ 范式更常见

Adversarial Training

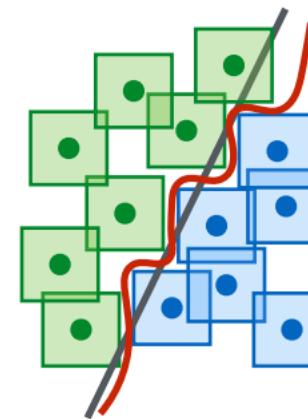
Adversarial training produces smooth decision boundary



正常边界



生成对抗样本



训练后

Early Adversarial Training Methods

- 2014年，Szegedy et al. 在其解释对抗样本的论文中已探索了对抗训练，用L-BFGS攻击对神经网络每一层生成对抗样本，并添加到训练过程中。
- **发现：深层对抗样本更能提高鲁棒性**
- 2015年，Goodfellow et al. 提出使用FGSM（单步）攻击生成的对抗样本来训练神经网络

$$\min_{\theta} \mathbb{E}_{(x,y) \in D} [\alpha \mathcal{L}_{\text{CE}}(f(\mathbf{x}), y) + (1 - \alpha) \mathcal{L}_{\text{CE}}(f(\mathbf{x}_{\text{adv}}), y)]$$
$$\mathbf{x}_{\text{adv}} = \mathbf{x} + \epsilon \cdot \text{sign}(\nabla_x \mathcal{L}_{\text{CE}}(f(\mathbf{x}), y))$$

Goodfellow等人并未使用中间层的对抗样本，因为发现中间层对抗样本没有提升



Min-max Robust Optimization

The First Proposal of Min-Max Optimization

$$\min_{\theta} \mathbb{E}_{(x,y) \in D} \max_{\|r\| \leq \epsilon} \mathcal{L}(f(\mathbf{x} + \mathbf{r}), y)$$

外部最小化
Outer minimization

内部最大化
Inner maximization

x_{adv}

Nokland et al. Improving back-propagation by adding an adversarial gradient. arXiv:1510.04189, 2015.

Huang et al. Learning with a strong adversary, ICLR 2016. Shaham et al. Understanding adversarial training: Increasing local stability of neural nets through robust optimization, arXiv:1511.05432, 2015

Virtual Adversarial Training

VAT: a method to improve generalization

$$\min_{\theta} \max_{\|x_{\text{adv}} - x\|_2 \leq \epsilon} \mathbb{E}_{(x,y) \in D} [\mathcal{L}_{\text{CE}}(f(x), y) + \lambda \mathcal{L}_{\text{KL}}(f(x), f(x_{\text{adv}}))]$$

- Differences to adversarial training
 - L2 regularized perturbation
 - Use both clean and adv examples for training
 - Use KL divergence to generate adv examples

Weaknesses of Early AT Methods

$$\min_{\theta} \mathbb{E}_{(\mathbf{x}, y) \in D} \max_{\|\mathbf{r}\| \leq \epsilon} \mathcal{L}(f(\mathbf{x} + \mathbf{r}), y)$$

- Use FGSM or BIM to solve the inner maximization problem
- FGSM and BIM were later found to be **weak attacks**
- Overfit to ϵ -robustness (**not robust to $<\epsilon$ attacks**)
- Overfit to single-step attacks (**not robust to multi-step attacks**)

✓ These methods are **fast!** Only takes **x2** time of standard training

PGD Adversarial Training

Defense	Dataset	Distance	Accuracy
Buckman et al. (2018)	CIFAR	0.031 (ℓ_∞)	0%*
Ma et al. (2018)	CIFAR	0.031 (ℓ_∞)	5%
Guo et al. (2018)	ImageNet	0.005 (ℓ_2)	0%*
Dhillon et al. (2018)	CIFAR	0.031 (ℓ_∞)	0%
Xie et al. (2018)	ImageNet	0.031 (ℓ_∞)	0%*
Song et al. (2018)	CIFAR	0.031 (ℓ_∞)	9%*
Samangouei et al. (2018)	MNIST	0.005 (ℓ_2)	55%**
Madry et al. (2018)	CIFAR	0.031 (ℓ_∞)	47%
Na et al. (2018)	CIFAR	0.015 (ℓ_∞)	15%

We got a survivor!



PGD Adversarial Training

A Saddle Point Problem

$$\min_{\theta} \mathbb{E}_{(x,y) \sim D} \max_{\|r\| \leq \epsilon} \mathcal{L}(f(x + r), y)$$

外部最小化
Outer minimization

内部最大化
Inner maximization

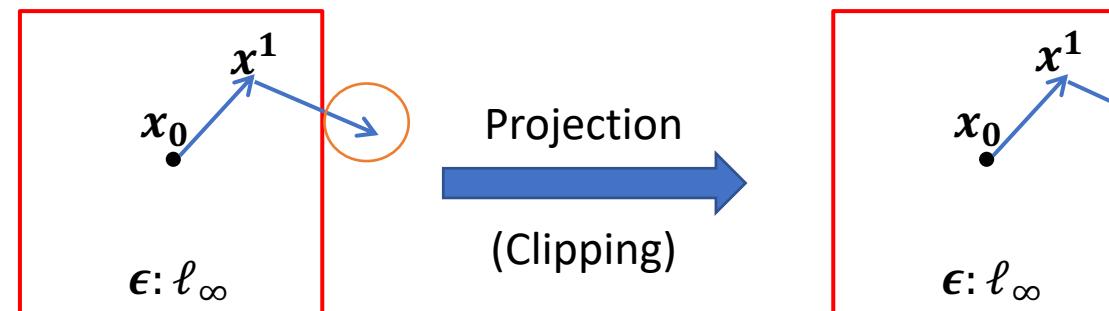
A saddle point (**constrained** bi-level optimization) problem

In constrained optimization, **Projected Gradient Descent (PGD)** is the best first-order solver

PGD Adversarial Training

Projected Gradient Descent (PGD)

$$\mathbf{x}_{\text{adv}}^{t+1} = \text{Proj}_{x+\mathcal{S}}(\mathbf{x}^t + \alpha \cdot \text{sign}(\nabla_{\mathbf{x}^t}) \mathcal{L}(\theta, \mathbf{x}^t, y))$$

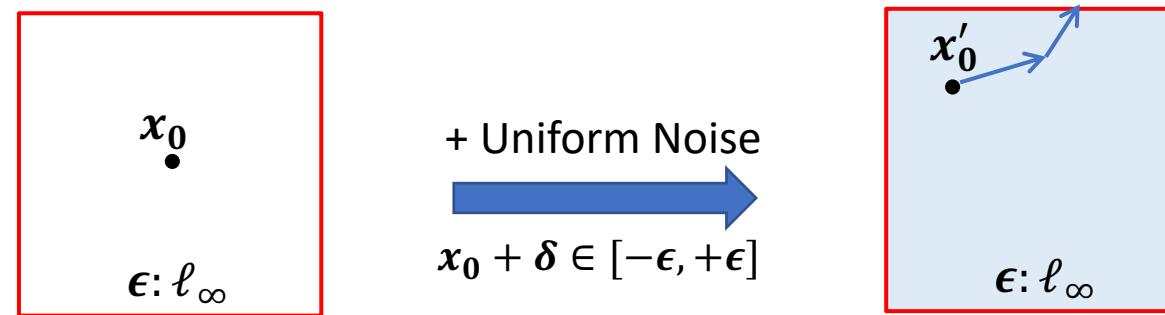


- PGD is an optimizer
- PGD is also known as an adv attack in the field of AML

PGD Adversarial Training

Projected Gradient Descent (PGD)

$$\mathbf{x}_{\text{adv}}^{t+1} = \text{Proj}_{x+\mathcal{S}}(\mathbf{x}^t + \alpha \cdot \text{sign}(\nabla_{\mathbf{x}^t}) \mathcal{L}(\theta, \mathbf{x}^t, y))$$



- Random initialization

PGD Adversarial Training

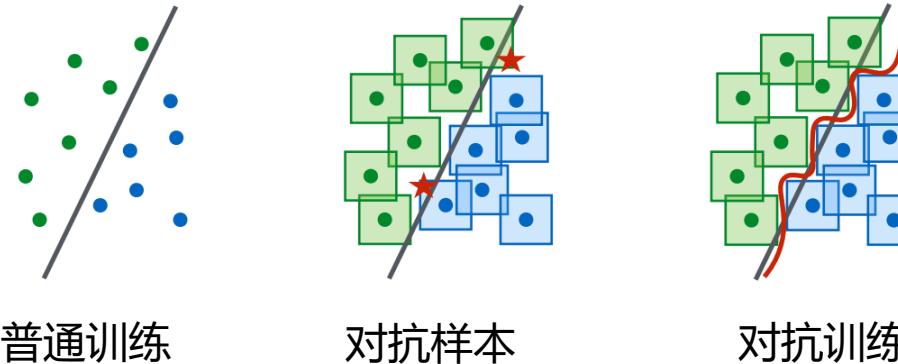
Characteristics of PGD adversarial training

- 只在对抗样本上训练模型
- 通用鲁棒性： $< \epsilon$ 鲁棒性和多步攻击鲁棒性
- 需要大容量模型
- 需要更多训练数据
- 对抗训练会产生平滑的决策边界
- 对抗训练会对内部激活产生一种截断效果
- 对抗训练会强制模型学习鲁棒特征
- 鲁棒性提升的同时干净准确率会下降
- 训练很耗时，相当于5-10倍普通训练

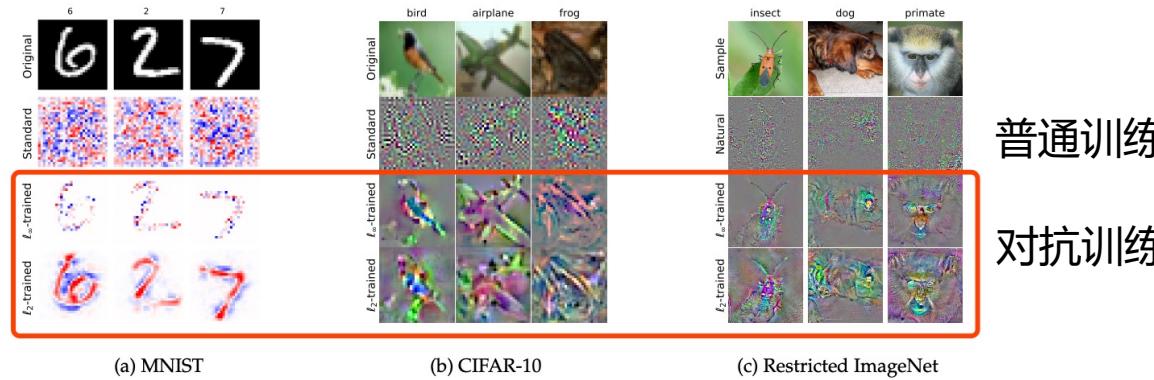


PGD Adversarial Training

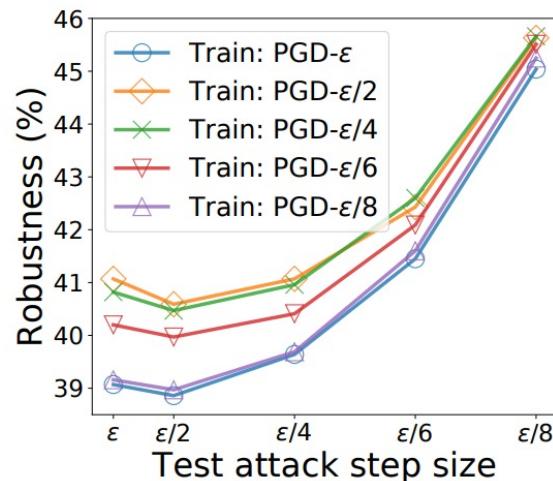
决策边界：



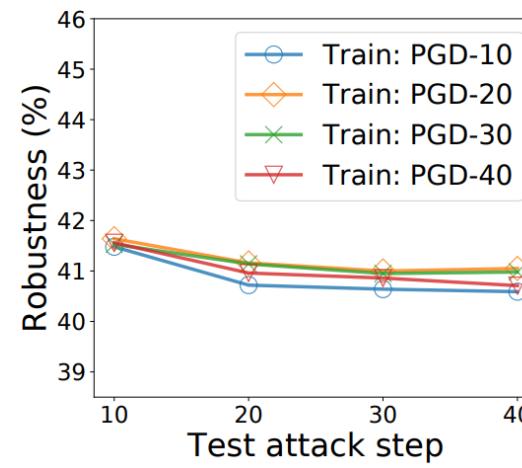
鲁棒特征：



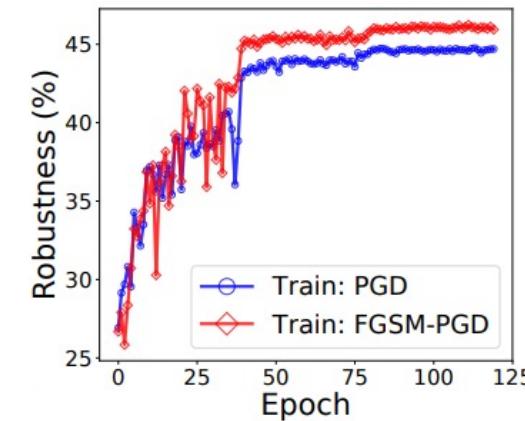
Dynamic Adversarial Training (DART)



PGD步长对鲁棒性影响



PGD步数对鲁棒性影响



训练初期使用简单攻击

- 对于10步对抗训练 (PGD-10 Adversarial Training) 来说:
 - 最优步长是 $\epsilon/2$ 和 $\epsilon/4$
 - 步数影响不大, 只要足够探索到 ϵ -ball的边界
- 训练初期使用弱对抗样本反而会提高鲁棒性

Dynamic Adversarial Training (DART)

How to measure the convergence of the inner maximization?

Definition (First-Order Stationary Condition (FOSC))

Given a sample $x^0 \in X$, let x^k be an intermediate example found at the k^{th} step of the inner maximization. The First-Order Stationary Condition of x^k is:

$$c(x^k) = \max_{x \in \chi} \langle x - x^k, \nabla_x f(\theta, x^k) \rangle,$$

where $X = \{x | \|x - x^0\|_\infty \leq \epsilon\}$ is the input domain of the ϵ -ball around normal example x^0 , $f(\theta, x^k) = \ell(h_\theta(x^k), y)$, and $\langle \cdot \rangle$ is the inner product.

FOSC:

- Inspired by the Frank-Wolfe gap for constrained min-max optimization.
- Smaller value of $c(x^k)$ indicates **better** convergence of the inner maximization.
- It has a close-form solution, affine invariant and norm independent.

Dynamic Adversarial Training (DART)

Dynamic Adversarial Training:

- Weak attack for early training, strong attack for later training
- **Weak attack improves generalization, strong attack improves final robustness.**

Convergence analysis:

Theorem 1. Suppose Assumptions 1, 2 and 3 hold. Let $\Delta = L_S(\theta^0) - \min_{\theta} L_S(\theta)$. If the step size of the outer minimization is set to $\eta_t = \eta = \min(1/L, \sqrt{\Delta/L\sigma^2T})$. Then the output of Algorithm 1 satisfies

$$\frac{1}{T} \sum_{t=0}^{T-1} \mathbb{E}[\|\nabla L_S(\theta^t)\|_2^2] \leq 4\sigma \sqrt{\frac{L\Delta}{T}} + \frac{5L_{\theta x}^2\delta}{\mu},$$

where $L = (L_{\theta x}L_{x\theta}/\mu + L_{\theta\theta})$.

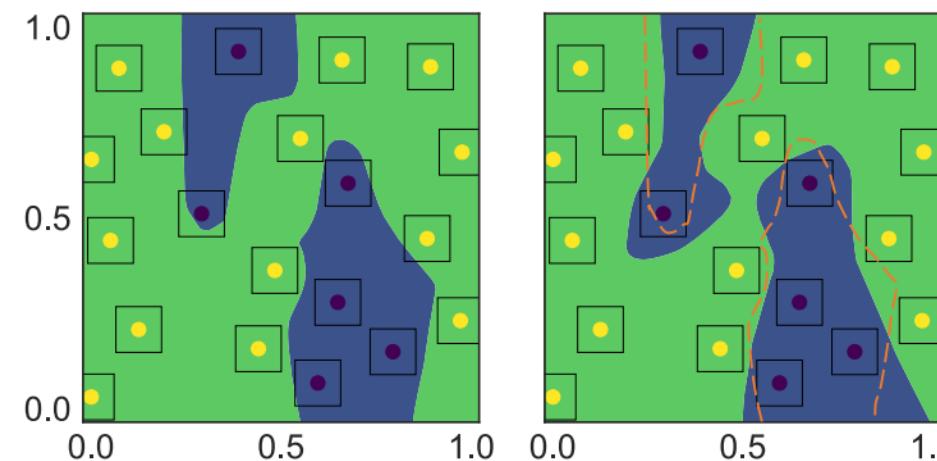
Robustness on CIFAR-10 with WideResNet

Defense	Clean	FGSM	PGD-20	C&W _∞
Madry's	87.3	56.1	45.8	46.8
Curriculum	77.43	57.17	46.06	42.28
Dynamic	85.03	63.53	48.70	47.27

DART improves robustness

Use distribution loss (KL) for inner and outer optimizations

$$\min_{\theta} \mathbb{E}_{(x,y) \sim D} \left[\underbrace{\mathcal{L}_{\text{CE}}(f(\mathbf{x}), y)}_{\text{提升准确率}} + \lambda \max_{\|\mathbf{r}\| \leq \epsilon} \underbrace{\mathcal{L}_{\text{KL}}(f(\mathbf{x}), f(\mathbf{x} + \mathbf{r}))}_{\text{提升鲁棒性}} \right]$$



Characteristics of TRADES

TRADES既改进了内部最大化又改进了外部最小化

- 使用KL监督对抗样本的生成，鲁棒性提升显著
- 干净样本也参与训练，有利于模型收敛和干净准确率
- 基于KL的对抗样本生成包含自适应的过程
- 能够训练得到比PGD对抗训练更平滑的决策边界

Winning solutions of
NeurIPS 2018
Adversarial Vision
Challenge



Experimental results of TRADES

Table 5. Comparisons of TRADES with prior defense models under white-box attacks.

Defense	Defense type	Under which attack	Dataset	Distance	$\mathcal{A}_{\text{nat}}(f)$	$\mathcal{A}_{\text{rob}}(f)$
Buckman et al. (2018)	gradient mask	Athalye et al. (2018)	CIFAR10	0.031 (ℓ_{∞})	-	0%
Ma et al. (2018)	gradient mask	Athalye et al. (2018)	CIFAR10	0.031 (ℓ_{∞})	-	5%
Dhillon et al. (2018)	gradient mask	Athalye et al. (2018)	CIFAR10	0.031 (ℓ_{∞})	-	0%
Song et al. (2018)	gradient mask	Athalye et al. (2018)	CIFAR10	0.031 (ℓ_{∞})	-	9%
Na et al. (2017)	gradient mask	Athalye et al. (2018)	CIFAR10	0.015 (ℓ_{∞})	-	15%
Wong et al. (2018)	robust opt.	FGSM ²⁰ (PGD)	CIFAR10	0.031 (ℓ_{∞})	27.07%	23.54%
Madry et al. (2018)	robust opt.	FGSM ²⁰ (PGD)	CIFAR10	0.031 (ℓ_{∞})	87.30%	47.04%
Zheng et al. (2016)	regularization	FGSM ²⁰ (PGD)	CIFAR10	0.031 (ℓ_{∞})	94.64%	0.15%
Kurakin et al. (2017)	regularization	FGSM ²⁰ (PGD)	CIFAR10	0.031 (ℓ_{∞})	85.25%	45.89%
Ross & Doshi-Velez (2017)	regularization	FGSM ²⁰ (PGD)	CIFAR10	0.031 (ℓ_{∞})	95.34%	0%
TRADES ($1/\lambda = 1.0$)	regularization	FGSM ²⁰ (PGD)	CIFAR10	0.031 (ℓ_{∞})	88.64%	49.14%
TRADES ($1/\lambda = 6.0$)	regularization	FGSM ²⁰ (PGD)	CIFAR10	0.031 (ℓ_{∞})	84.92%	56.61%
TRADES ($1/\lambda = 1.0$)	regularization	DeepFool (ℓ_{∞})	CIFAR10	0.031 (ℓ_{∞})	88.64%	59.10%
TRADES ($1/\lambda = 6.0$)	regularization	DeepFool (ℓ_{∞})	CIFAR10	0.031 (ℓ_{∞})	84.92%	61.38%
TRADES ($1/\lambda = 1.0$)	regularization	LBFSGSAAttack	CIFAR10	0.031 (ℓ_{∞})	88.64%	84.41%
TRADES ($1/\lambda = 6.0$)	regularization	LBFSGSAAttack	CIFAR10	0.031 (ℓ_{∞})	84.92%	81.58%
TRADES ($1/\lambda = 1.0$)	regularization	MI-FGSM	CIFAR10	0.031 (ℓ_{∞})	88.64%	51.26%
TRADES ($1/\lambda = 6.0$)	regularization	MI-FGSM	CIFAR10	0.031 (ℓ_{∞})	84.92%	57.95%
TRADES ($1/\lambda = 1.0$)	regularization	C&W	CIFAR10	0.031 (ℓ_{∞})	88.64%	84.03%
TRADES ($1/\lambda = 6.0$)	regularization	C&W	CIFAR10	0.031 (ℓ_{∞})	84.92%	81.24%
Samangouei et al. (2018)	gradient mask	Athalye et al. (2018)	MNIST	0.005 (ℓ_2)	-	55%
Madry et al. (2018)	robust opt.	FGSM ⁴⁰ (PGD)	MNIST	0.3 (ℓ_{∞})	99.36%	96.01%
TRADES ($1/\lambda = 6.0$)	regularization	FGSM ⁴⁰ (PGD)	MNIST	0.3 (ℓ_{∞})	99.48%	96.07%
TRADES ($1/\lambda = 6.0$)	regularization	C&W	MNIST	0.005 (ℓ_2)	99.48%	99.46%

TRADES vs VAT vs ALP

TRADES:

$$\min_{\theta} \mathbb{E}_{(x,y) \sim D} \left[\underbrace{\mathcal{L}_{\text{CE}}(f(\mathbf{x}), y)}_{\text{提升准确率}} + \lambda \underbrace{\max_{\|\mathbf{r}\| \leq \epsilon} \mathcal{L}_{\text{KL}}(f(\mathbf{x}), f(\mathbf{x} + \mathbf{r}))}_{\text{提升鲁棒性}} \right]$$

Virtual Adversarial Training:

$$\min_{\theta} \max_{\|\mathbf{x}_{\text{adv}} - \mathbf{x}\|_2 \leq \epsilon} \mathbb{E}_{(x,y) \in D} [\mathcal{L}_{\text{CE}}(f(\mathbf{x}), y) + \lambda \mathcal{L}_{\text{KL}}(f(\mathbf{x}), f(\mathbf{x}_{\text{adv}}))]$$

Adversarial Logits Pairing:

$$\min_{\theta} \mathbb{E}_{(x,y) \in D} \left[\max_{\|\mathbf{r}\| \leq \epsilon} \mathcal{L}_{\text{CE}}(f(\mathbf{x} + \mathbf{r}), y) + \lambda \|f(\mathbf{x} + \mathbf{r}) - f(\mathbf{x})\|_2^2 \right]$$

相似的优化框架，不同的损失选择，结果差异很大

Zhang et al. "Theoretically principled trade-off between robustness and accuracy." ICML, 2019.

Miyato et al. Distributional smoothing with virtual adversarial training. ICLR 2016.

Kannan, Harini, Alexey Kurakin, and Ian Goodfellow. "Adversarial logit pairing." *arXiv preprint arXiv:1803.06373* (2018).



MART: Misclassification Aware adveRsarial Training

Min-max Adversarial Training:

$$\min_{\theta} \frac{1}{n} \sum_{i=1}^n \max_{\|x_i - x_i^0\| \leq \epsilon} \mathcal{L}(f_{\theta}(x_i), y_i)$$

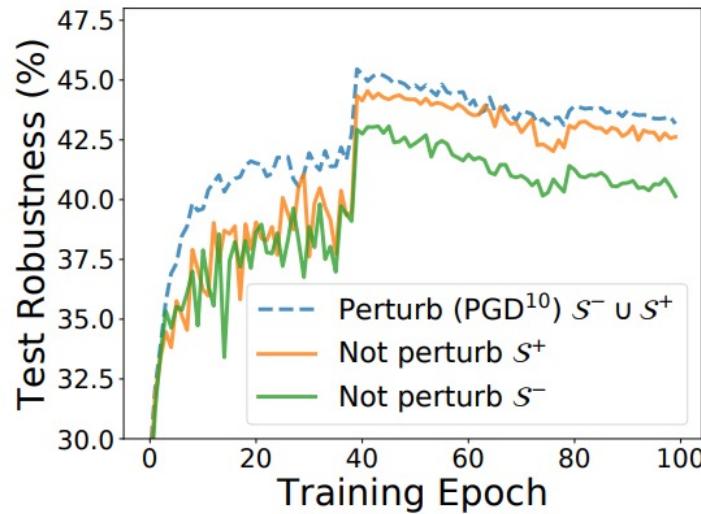
where, x_i^0 is a natural (clean) training sample, y_i is the label of x_i^0 .

Adversarial examples are only defined on **correctly classified** examples, what about **misclassified examples** ?

MART: Misclassification Aware adveRsarial Training

The influence of misclassified and correctly classified examples:

- A pre-trained network to select the **same size (13%)**
 - Subset of misclassified examples S^-
 - Subset of correctly classified examples S^+

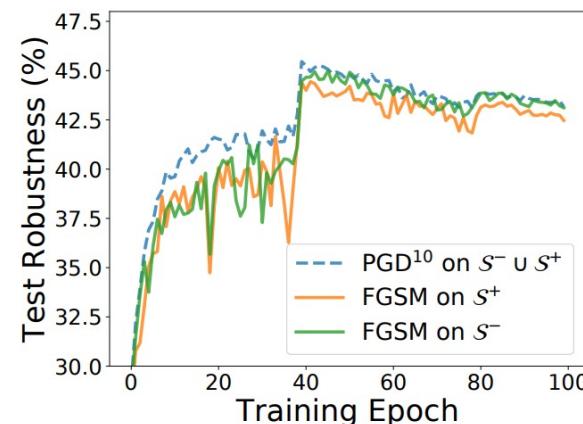


Misclassified examples have a **significant impact** on the final robustness!

MART: Misclassification Aware adveRsarial Training

- For **inner maximization** process:

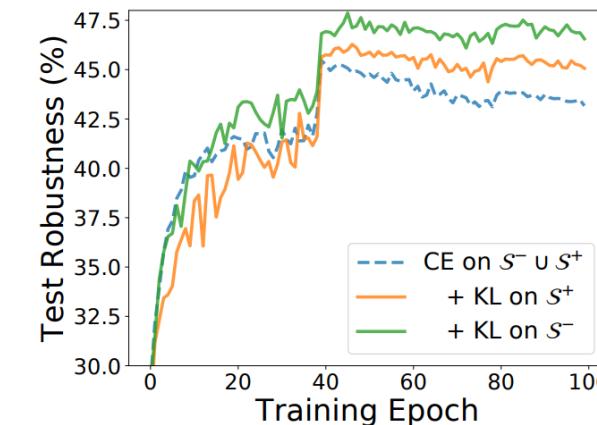
- Weak attack on misclassified examples S^-
- Weak attack on correctly classified examples S^+



(b) Inner maximization

- For **outer minimization** process:

- Regularization on misclassified examples S^-
- Regularization on correctly classified examples S^+



(c) Outer minimization

different maximization techniques have
negligible effect

different minimization techniques have
significant effect

Wang, et al. "Improving adversarial robustness requires revisiting misclassified examples." ICLR, 2019.

MART: Misclassification Aware adveRsarial Training

Misclassification aware adversarial risk:

- Adversarial risk:

$$\mathcal{R}(h_{\theta}) = \frac{1}{n} \sum_{i=1}^n \max_{\mathbf{x}'_i \in \mathcal{B}_{\epsilon}(\mathbf{x}_i)} \mathbb{1}(h_{\theta}(\mathbf{x}'_i) \neq y_i),$$

- Correctly classified and misclassified example:

$$\mathcal{S}_{h_{\theta}}^{+} = \{i : i \in [n], h_{\theta}(\mathbf{x}_i) = y_i\} \quad \text{and} \quad \mathcal{S}_{h_{\theta}}^{-} = \{i : i \in [n], h_{\theta}(\mathbf{x}_i) \neq y_i\}$$

- Misclassification aware adversarial risk:

$$\begin{aligned} \min_{\theta} \mathcal{R}_{\text{misc}}(h_{\theta}) &:= \frac{1}{n} \left(\sum_{i \in \mathcal{S}_{h_{\theta}}^{+}} \mathcal{R}^{+}(h_{\theta}, \mathbf{x}_i) + \sum_{i \in \mathcal{S}_{h_{\theta}}^{-}} \mathcal{R}^{-}(h_{\theta}, \mathbf{x}_i) \right) \\ &= \frac{1}{n} \sum_{i=1}^n \{ \mathbb{1}(h_{\theta}(\hat{\mathbf{x}}'_i) \neq y_i) + \mathbb{1}(h_{\theta}(\mathbf{x}_i) \neq h_{\theta}(\hat{\mathbf{x}}'_i)) \cdot \mathbb{1}(h_{\theta}(\mathbf{x}_i) \neq y_i) \} \end{aligned}$$

MART: Misclassification Aware adveRsarial Training

- Surrogate loss functions (existing methods and MART):

Defense Method	Loss Function
<i>Standard</i>	$\text{CE}(\mathbf{p}(\hat{\mathbf{x}}', \boldsymbol{\theta}), y)$
ALP	$\text{CE}(\mathbf{p}(\hat{\mathbf{x}}', \boldsymbol{\theta}), y) + \lambda \cdot \ \mathbf{p}(\hat{\mathbf{x}}', \boldsymbol{\theta}) - \mathbf{p}(\mathbf{x}, \boldsymbol{\theta})\ _2^2$
CLP	$\text{CE}(\mathbf{p}(\mathbf{x}, \boldsymbol{\theta}), y) + \lambda \cdot \ \mathbf{p}(\hat{\mathbf{x}}', \boldsymbol{\theta}) - \mathbf{p}(\mathbf{x}, \boldsymbol{\theta})\ _2^2$
TRADES	$\text{CE}(\mathbf{p}(\mathbf{x}, \boldsymbol{\theta}), y) + \lambda \cdot \text{KL}(\mathbf{p}(\mathbf{x}, \boldsymbol{\theta}) \parallel \mathbf{p}(\hat{\mathbf{x}}', \boldsymbol{\theta}))$
MMA	$\text{CE}(\mathbf{p}(\hat{\mathbf{x}}', \boldsymbol{\theta}), y) \cdot \mathbb{1}(h_{\boldsymbol{\theta}}(\mathbf{x}) = y) + \text{CE}(\mathbf{p}(\mathbf{x}, \boldsymbol{\theta}), y) \cdot \mathbb{1}(h_{\boldsymbol{\theta}}(\mathbf{x}) \neq y)$
MART	$\text{BCE}(\mathbf{p}(\hat{\mathbf{x}}', \boldsymbol{\theta}), y) + \lambda \cdot \text{KL}(\mathbf{p}(\mathbf{x}, \boldsymbol{\theta}) \parallel \mathbf{p}(\hat{\mathbf{x}}', \boldsymbol{\theta})) \cdot (1 - \mathbf{p}_y(\mathbf{x}, \boldsymbol{\theta}))$

- Semi-supervised extension of MART: $\text{BCE}(\mathbf{p}(\hat{\mathbf{x}}'_i, \boldsymbol{\theta}), y_i) = -\log(\mathbf{p}_{y_i}(\hat{\mathbf{x}}'_i, \boldsymbol{\theta})) - \log(1 - \max_{k \neq y_i} \mathbf{p}_k(\hat{\mathbf{x}}'_i, \boldsymbol{\theta}))$

$$\mathcal{L}_{\text{semi}}^{\text{MART}}(\boldsymbol{\theta}) = \sum_{i \in \mathcal{S}_{\text{sup}}} \ell_{\text{sup}}^{\text{MART}}(\mathbf{x}_i, y_i; \boldsymbol{\theta}) + \gamma \cdot \sum_{i \in \mathcal{S}_{\text{unsup}}} \ell_{\text{unsup}}^{\text{MART}}(\mathbf{x}_i, y_i; \boldsymbol{\theta})$$

MART: Misclassification Aware adveRsarial Training

Robustness of MART:

- White-box robustness: **ResNet-18**, CIFAR-10, $\epsilon=8/255$

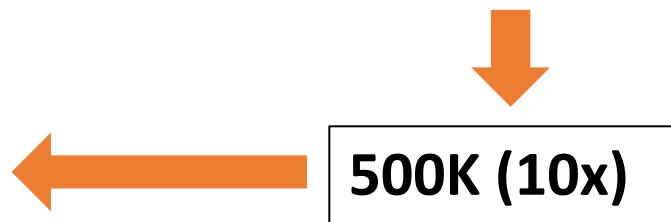
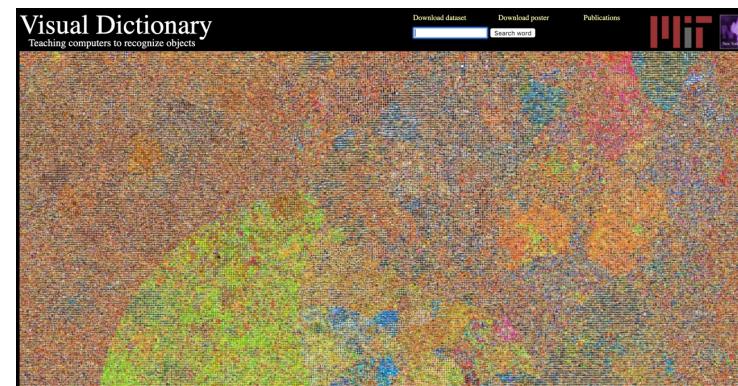
Defense	MNIST				CIFAR-10			
	Natural	FGSM	PGD ²⁰	CW _{∞}	Natural	FGSM	PGD ²⁰	CW _{∞}
<i>Standard</i>	99.11	97.17	94.62	94.25	84.44	61.89	47.55	45.98
MMA	98.92	97.25	95.25	94.77	84.76	62.08	48.33	45.77
Dynamic	98.96	97.34	95.27	94.85	83.33	62.47	49.40	46.94
TRADES	99.25	96.67	94.58	94.03	82.90	62.82	50.25	48.29
MART	98.74	97.87	96.48	96.10	83.07	65.65	55.57	54.87

- White-box robustness: **WideResNet-34-10**, CIFAR-10, $\epsilon=8/255$

Defense	Natural	FGSM		PGD ²⁰		PGD ¹⁰⁰		CW _{∞}	
		Best	Last	Best	Last	Best	Last	Best	Last
<i>Standard</i>	87.30	56.10	56.10	52.68	49.31	51.55	49.03	50.73	48.47
Dynamic	84.51	63.53	63.53	55.03	51.70	54.12	50.07	51.34	49.27
TRADES	84.22	64.70	64.70	56.40	53.16	55.68	51.27	51.98	51.12
MART	84.17	67.51	67.51	58.56	57.39	57.88	55.04	54.58	54.53

Using More Data to Improve Robustness

80 Million Tiny Images



Select (carefully) 100K/500K images into CIAFR-10



Alayrac et al. "Are labels required for improving adversarial robustness?." *NeurIPS*, 2019.
Carmon et al. "Unlabeled data improves adversarial robustness." *NeurIPS* 2019

UAT & RST

Unsupervised Adversarial Training (UAT):

$$\mathcal{L}(\theta) = \mathcal{L}_{\text{sup}}(\theta) + \lambda \mathcal{L}_{\text{unsup}}(\theta)$$

$$\left[\begin{array}{l} \mathcal{L}_{\text{unsup}}^{OT}(\theta) = \mathbb{E}_{x \sim P(X)} \sup_{x' \in \mathcal{N}_\epsilon(x)} \mathcal{D}(p_{\hat{\theta}}(.|x), p_\theta(.|x')) \\ \mathcal{L}_{\text{unsup}}^{FT}(\theta) = \mathbb{E}_{x \sim P(X)} \sup_{x' \in \mathcal{N}_\epsilon(x)} \text{xent}(\hat{y}(x), p_\theta(.|x')) \end{array} \right]$$

Robust Self-Training (RST):

$$\text{minimizing } \sum_{i=1}^n L_{\text{robust}}(\theta, x_i, y_i) + w \sum_{i=1}^{\tilde{n}} L_{\text{robust}}(\theta, \tilde{x}_i, \tilde{y}_i) \rightarrow \left[\begin{array}{l} \tilde{x}_i: \text{new samples} \\ \tilde{y}_i : \text{pseudo-labels} \end{array} \right]$$

Alayrac et al. "Are labels required for improving adversarial robustness?." *NeurIPS*, 2019.

Carmon et al. "Unlabeled data improves adversarial robustness." *NeurIPS* 2019

State-of-the-art: AT Methods

ROBUSTBENCH

Leaderboards Paper FAQ Contribute Model Zoo 🚀

Leaderboard: CIFAR-10, $\ell_\infty = 8/255$, untargeted attack

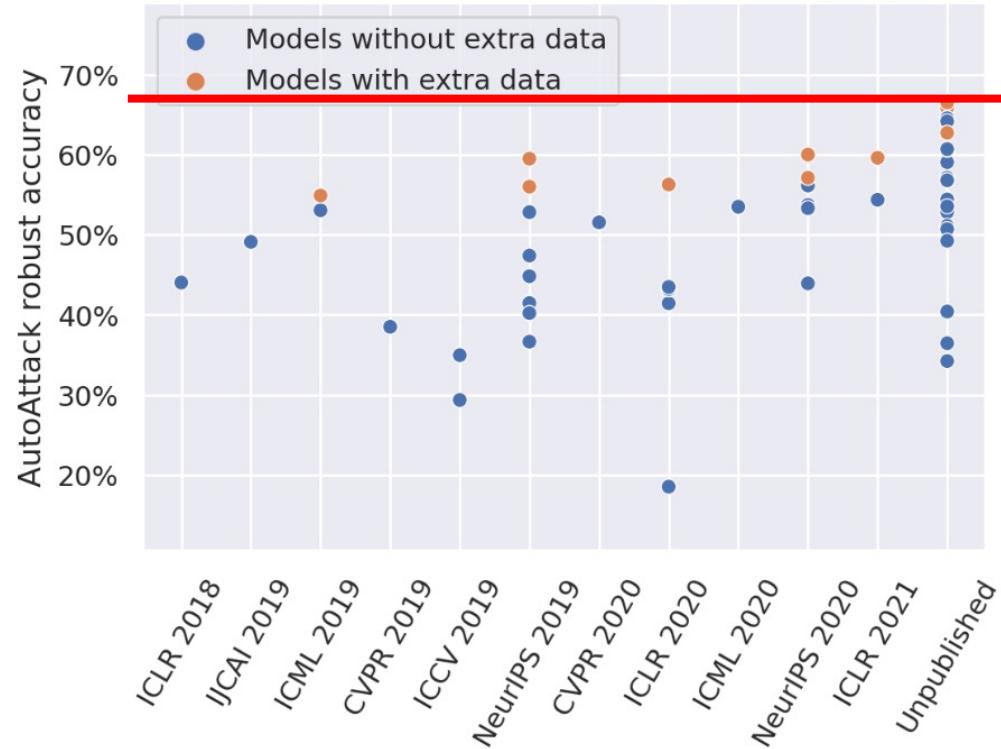
Show 15 entries Search: Papers, architectures, ve

Rank	Method	Standard accuracy	AutoAttack robust accuracy	Best known robust accuracy	AA eval. potentially unreliable	Extra data	Architecture	Venue
1	Robust Principles: Architectural Design Principles for Adversarially Robust CNNs <i>It uses additional 50M synthetic images in training.</i>	93.27%	71.07%	71.07%	×	×	RaWideResNet-70-16	BMVC 2023
2	Better Diffusion Models Further Improve Adversarial Training <i>It uses additional 50M synthetic images in training.</i>	93.25%	70.69%	70.69%	×	×	WideResNet-70-16	ICML 2023
3	Improving the Accuracy-Robustness Trade-off of Classifiers via Adaptive Smoothing <i>It uses an ensemble of networks. The robust base classifier uses 50M synthetic images.</i>	95.23%	68.06%	68.06%	×	✓	ResNet-152 + WideResNet-70-16 + mixing network	arXiv, Jan 2023
4	Decoupled Kullback-Leibler Divergence Loss <i>It uses additional 20M synthetic images in training.</i>	92.16%	67.73%	67.73%	×	×	WideResNet-28-10	arXiv, May 2023
5	Better Diffusion Models Further Improve Adversarial Training <i>It uses additional 20M synthetic images in training.</i>	92.44%	67.31%	67.31%	×	×	WideResNet-28-10	ICML 2023
6	Fixing Data Augmentation to Improve Adversarial Robustness <i>66.56% robust accuracy is due to the original evaluation (AutoAttack + MultiTargeted)</i>	92.23%	66.58%	66.56%	×	✓	WideResNet-70-16	arXiv, Mar 2021
7	Improving Robustness using Generated Data <i>It uses additional 100M synthetic images in training. 66.10% robust accuracy is due to the original evaluation (AutoAttack + MultiTargeted)</i>	88.74%	66.11%	66.10%	×	×	WideResNet-70-16	NeurIPS 2021

数据，数据，还是数据！！

- 数据增广
- 数据生成

State-of-the-art: AT Methods



Clean accuracy: 94% vs 66% (robustness)

DNN: WideResNet-70-16; Dataset: CIFAR-10
Perturbation: $\epsilon = 8/255$;
Evaluation attack: AutoAttack

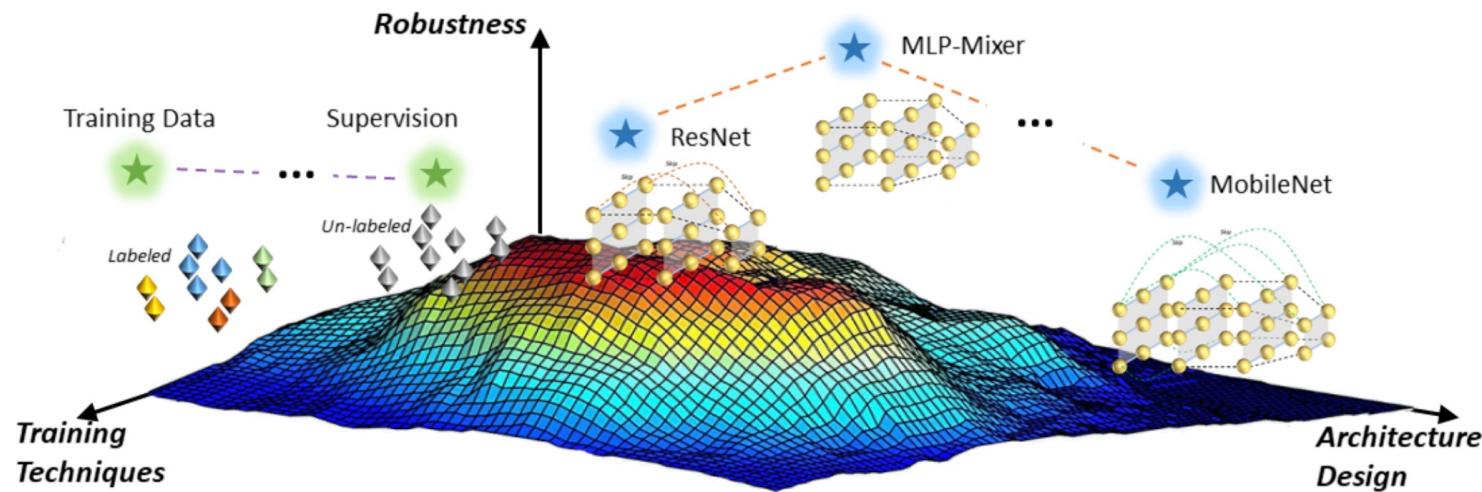
State-of-the-art: DNN Architecture



Home LeaderBoards API Docs Model Zoo Contact Toolkit Paper

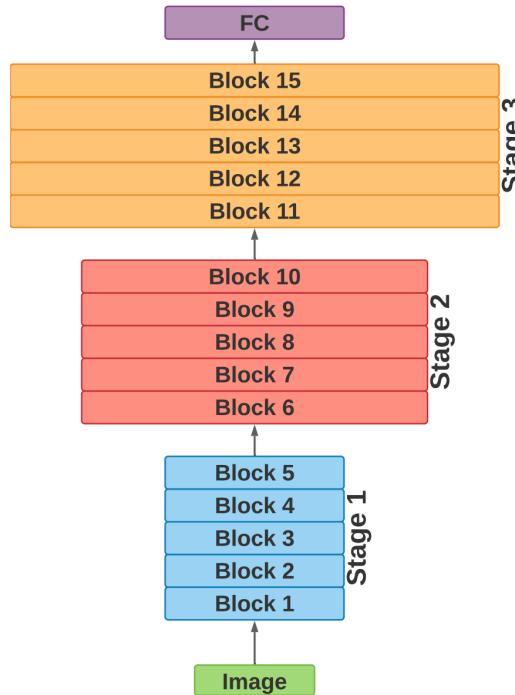
RobustART

RobustART is the first comprehensive **Robust**ness investigation benchmark on large-scale dataset ImageNet regarding **AR**chitectural design (**49** human-designed off-the-shelf architectures and **1200+** neural architecture searched networks) and **T**raining techniques (**10+** general ones e.g., extra training data, etc) towards diverse noises (adversarial, natural, and system noises). Our benchmark (including open-source toolkit, pre-trained model zoo, datasets, and analyses): (1) presents an open-source platform for conducting comprehensive evaluation on diverse robustness types; (2) provides a variety of pre-trained models with different training techniques to facilitate robustness evaluation; (3) proposes a new view to better understand the mechanism towards designing robust DNN architectures, backed up by the analysis. We will continuously contribute to building this ecosystem for the community.

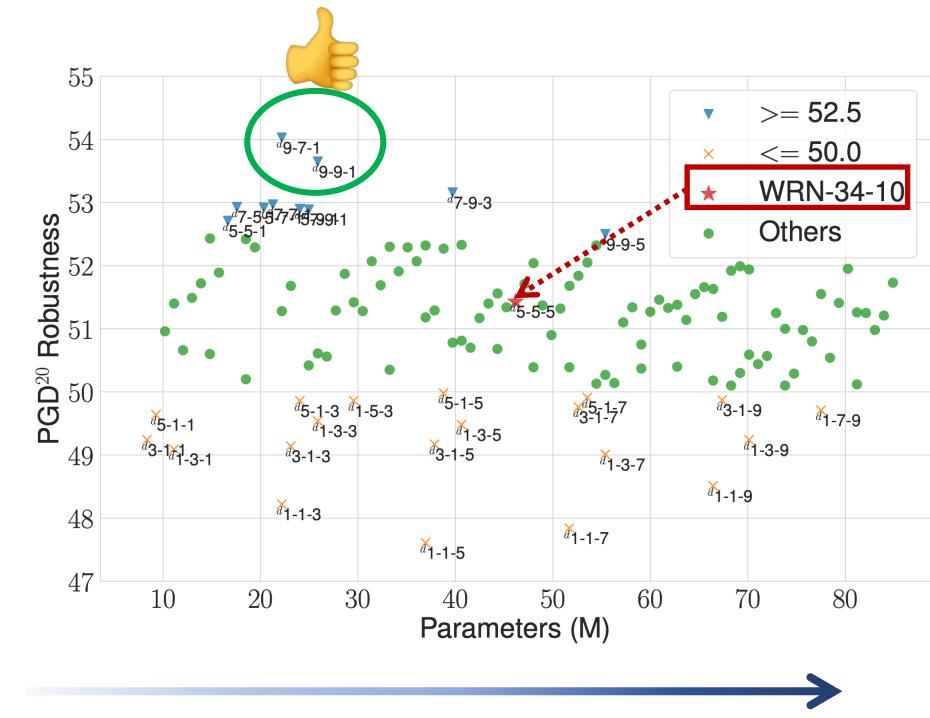


State-of-the-art: DNN Architecture

Reduce the deep layers can improve robustness



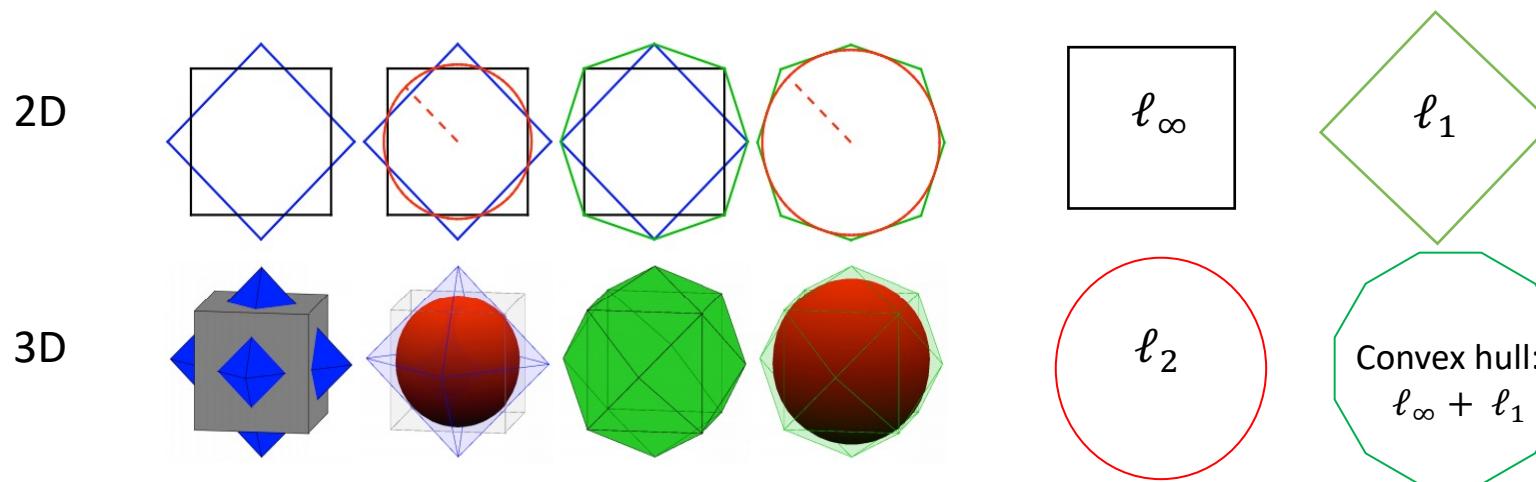
WideResNet-34-10



Depth Grid Search

Certified Defense vs Empirical Defense

- **Certified robustness** (Sinha et al. 2018; Cohen et al. 2018; Lee et al. 2019)
 - Gaussian randomized smoothing -> robustness with the ℓ_2 norm ball
 - Laplacian randomized smoothing -> ℓ_1 robustness
 - Uniform randomized smoothing -> ℓ_0 robustness
 - **Pros:** robustness guarantees, **严格的鲁棒性下界证明**
 - **Cons:** 1) deep networks are hard to certify, 2) guarantees are loose, 3) need to train the model in certain ways



关键点：如何在神经网络里传播边界？

Existing Challenges

- How to attack large language/vision/multi-model models
- How to defend large language/vision/multi-model models
- How to adapt adv training for different applications
- How to reduce the cost of defense: acceleration, loss of clean acc
- How to combine adv detection with adversarial training
- How to include adv training into the pretraining/finetuning pipeline



谢谢 !

